

# Machine Learning Fundamentals- DTSC102

## Lecture 2 Unsupervised Learning: Clustering

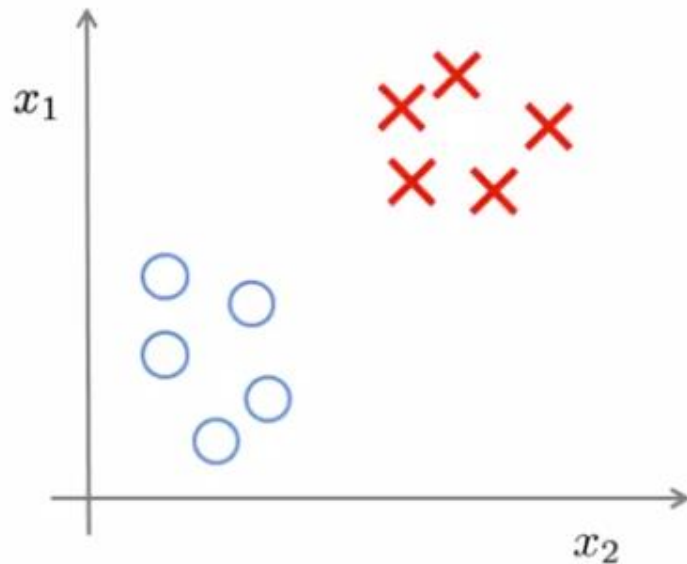
Course Instructors: Dr.-Ing. Maggie Mashaly  
maggie.ezzat@guc.edu.eg  
C3.220

# Contents

- Unsupervised learning
- Clustering
- K-means Algorithm
- Examples
- Advantages & Disadvantages
- Implementation Details

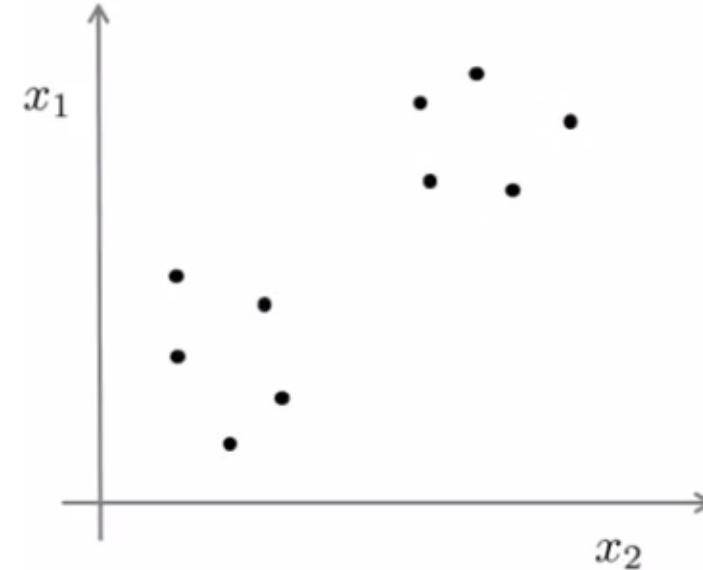
# Supervised vs. Unsupervised Learning

## Supervised Learning



**Training set:**  
 $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), (x^{(3)}, y^{(3)}), \dots, (x^{(m)}, y^{(m)})\}$

## Unsupervised Learning



**Training set:**  
 $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots, x^{(m)}\}$

# Supervised Learning: Classification

Training set of labeled docs



SPORTS



WORLD NEWS

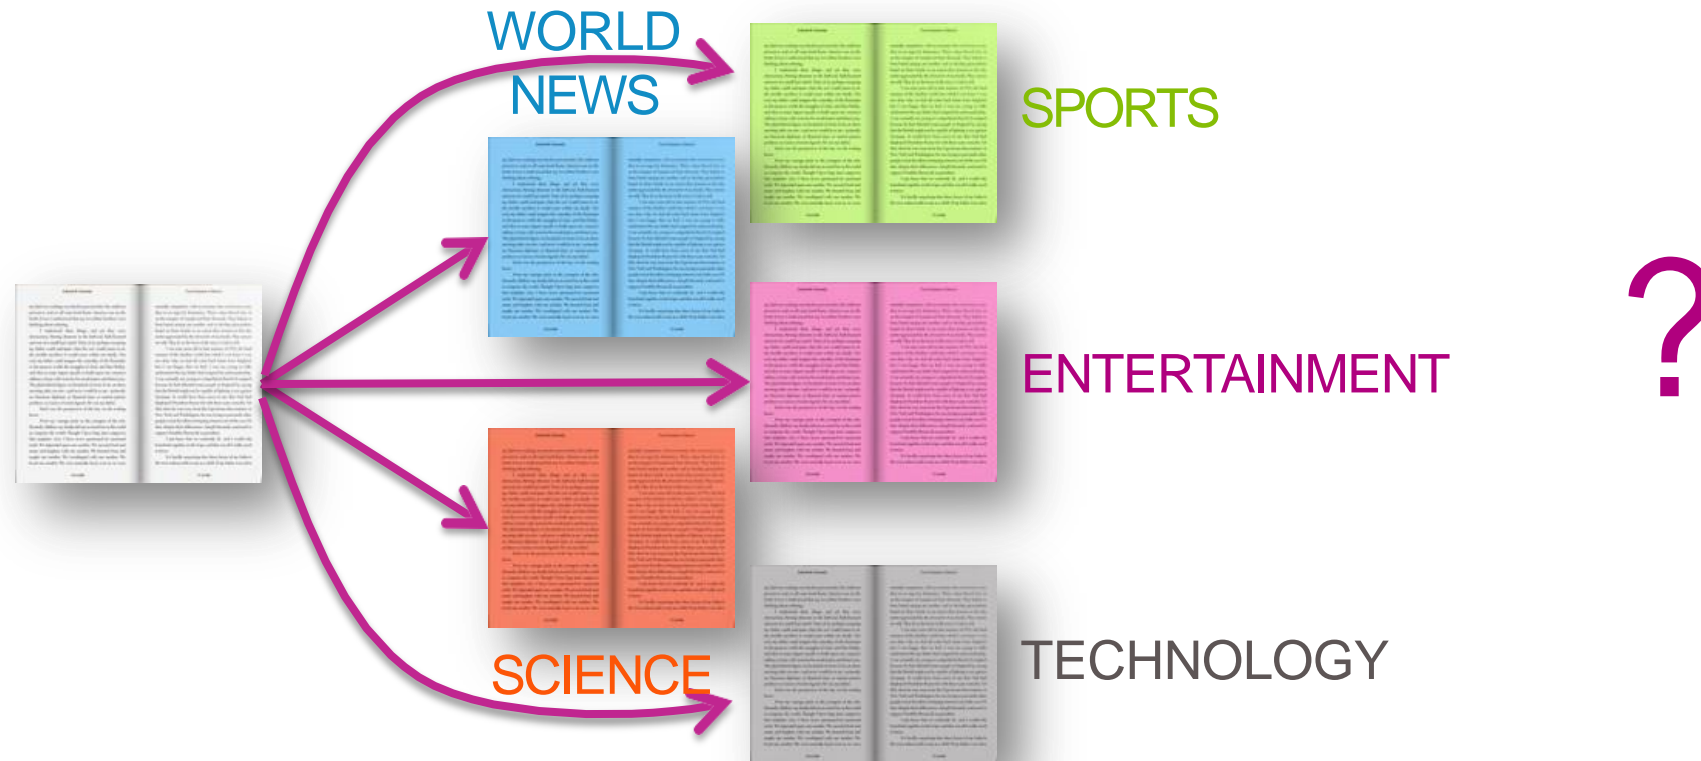


ENTERTAINMENT



SCIENCE

# Supervised Learning: Multi-class Classification



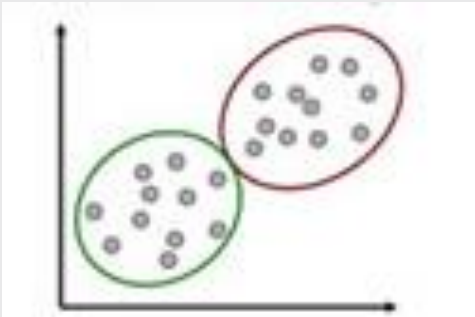
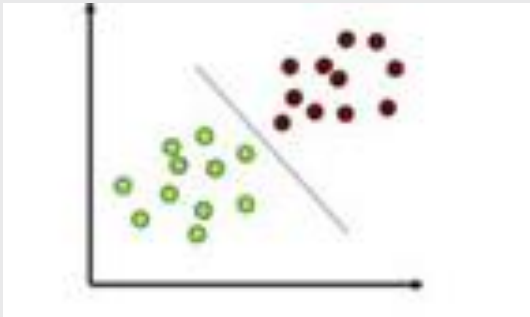
But what if labels were not provided...?

# Unsupervised Learning: Clustering

➤ **Goal:**

Finding structure within the data, usually by dividing it into **Clusters**

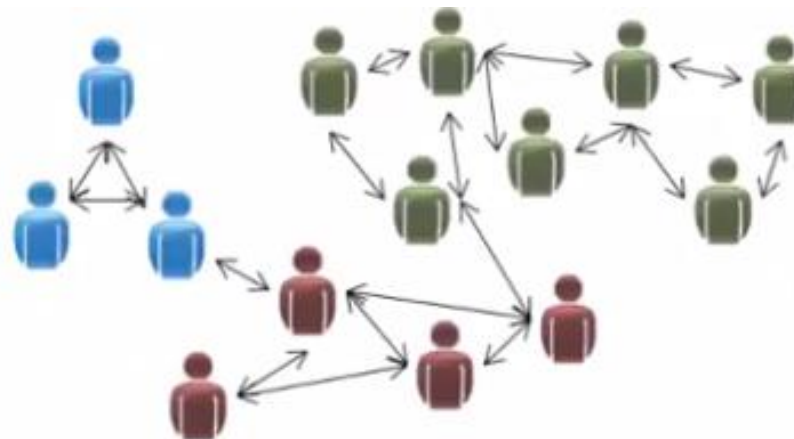
➤ **But mind the difference:**

Clustering	Classification
<ul style="list-style-type: none"><li>• Data is not labeled</li><li>• Group points that are “close” to each other</li><li>• Identify structure or patterns in data</li><li>• Unsupervised learning</li></ul>	<ul style="list-style-type: none"><li>• Labeled data points</li><li>• Want a “rule” that assigns labels to new points</li><li>• Supervised learning</li></ul>
	

# Clustering: Use cases



Market segmentation



Social network analysis



Organize computing clusters



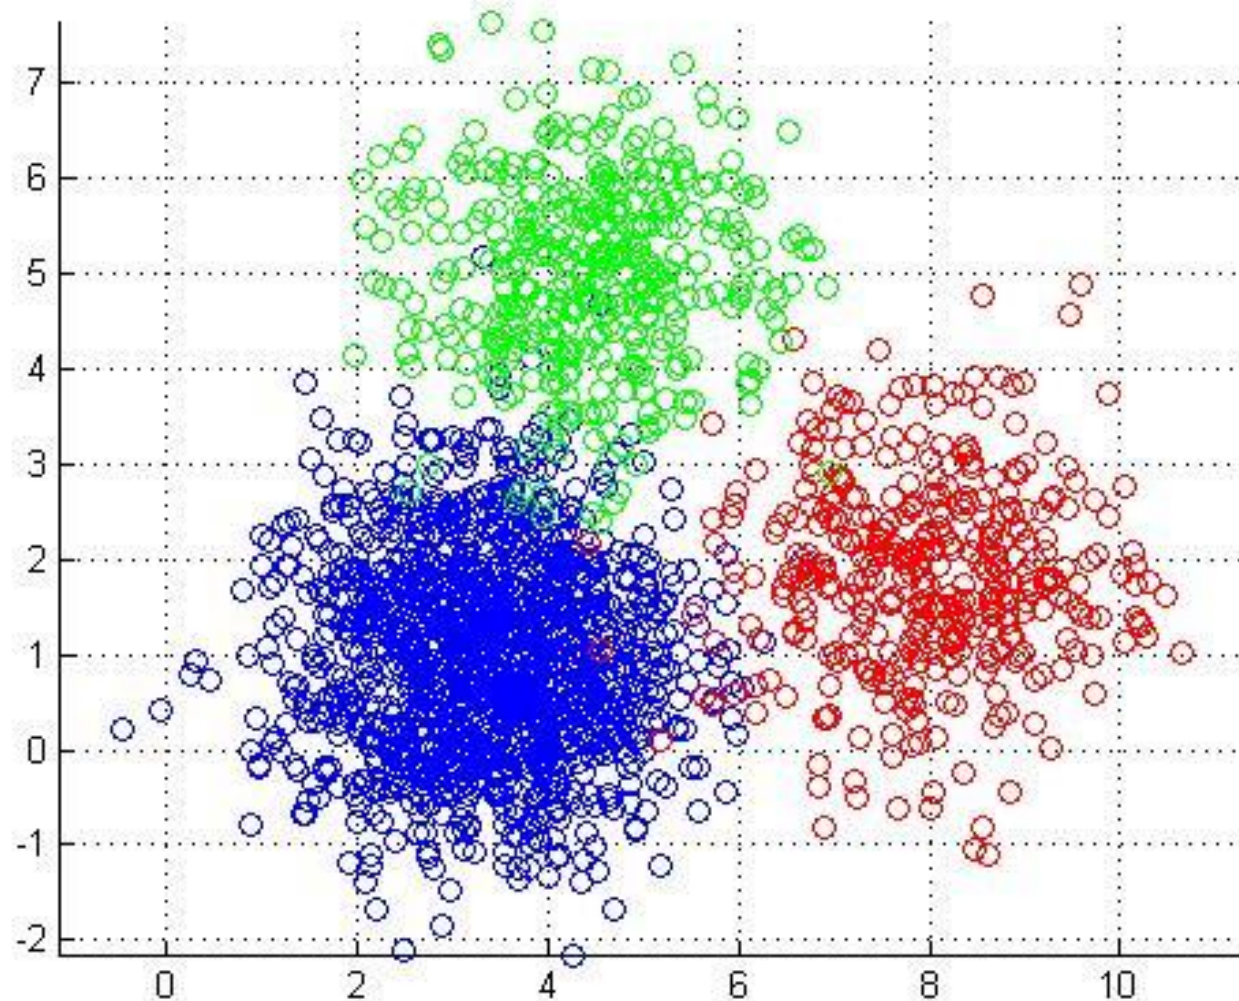
Astronomical data analysis

# Clustering: Use cases

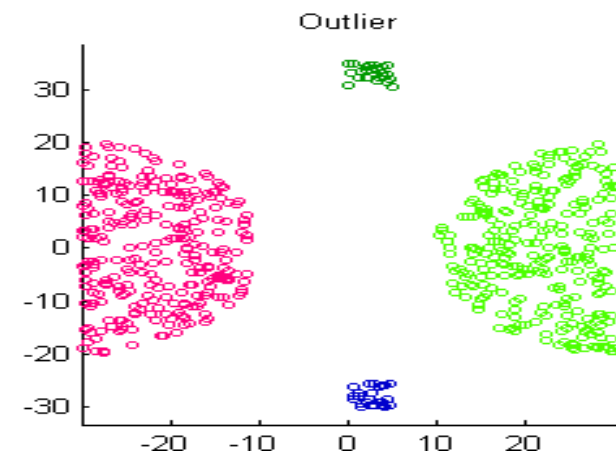
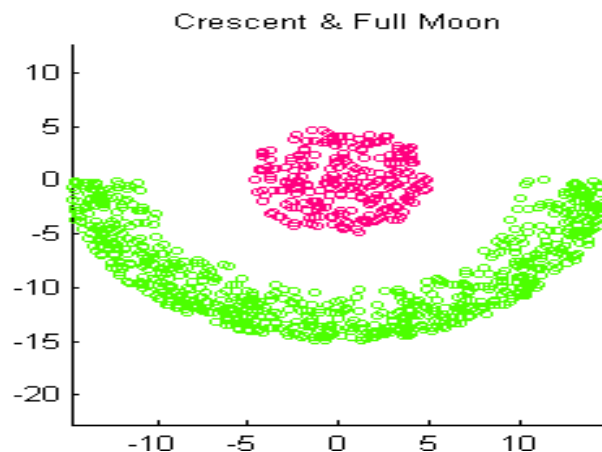
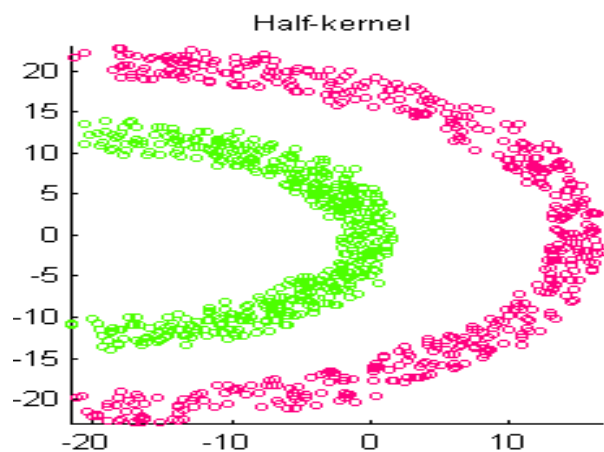
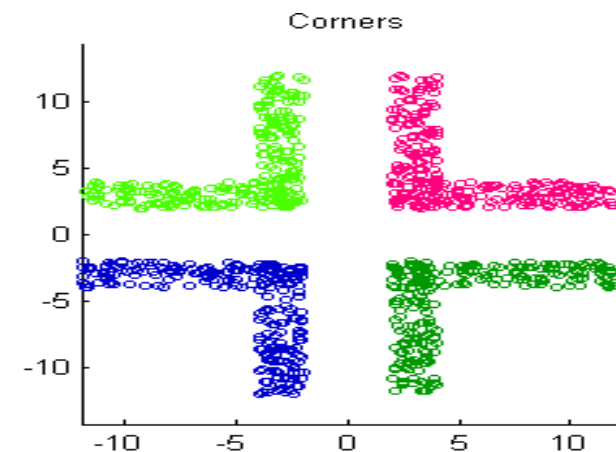
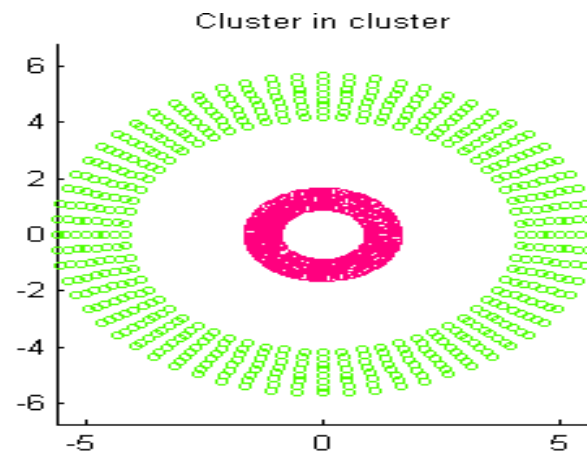
- Data summarization, compression, and reduction
  - *Examples: Image processing or vector quantization*
- Collaborative filtering, recommendation systems, or customer segmentation
  - *Finding like-minded users or similar products*
- Dynamic trend detection
  - *Clustering stream data and detecting trends and patterns*
- Multimedia data analysis, biological data analysis, and social network analysis
  - *Example: Clustering images or video/audio clips, gene/protein sequences, etc.*
- A key intermediate step for other data mining tasks
  - *Generating a compact summary of data for classification, pattern discovery, and hypothesis generation and testing*
- Outlier detection: Outliers - those “far away” from any cluster



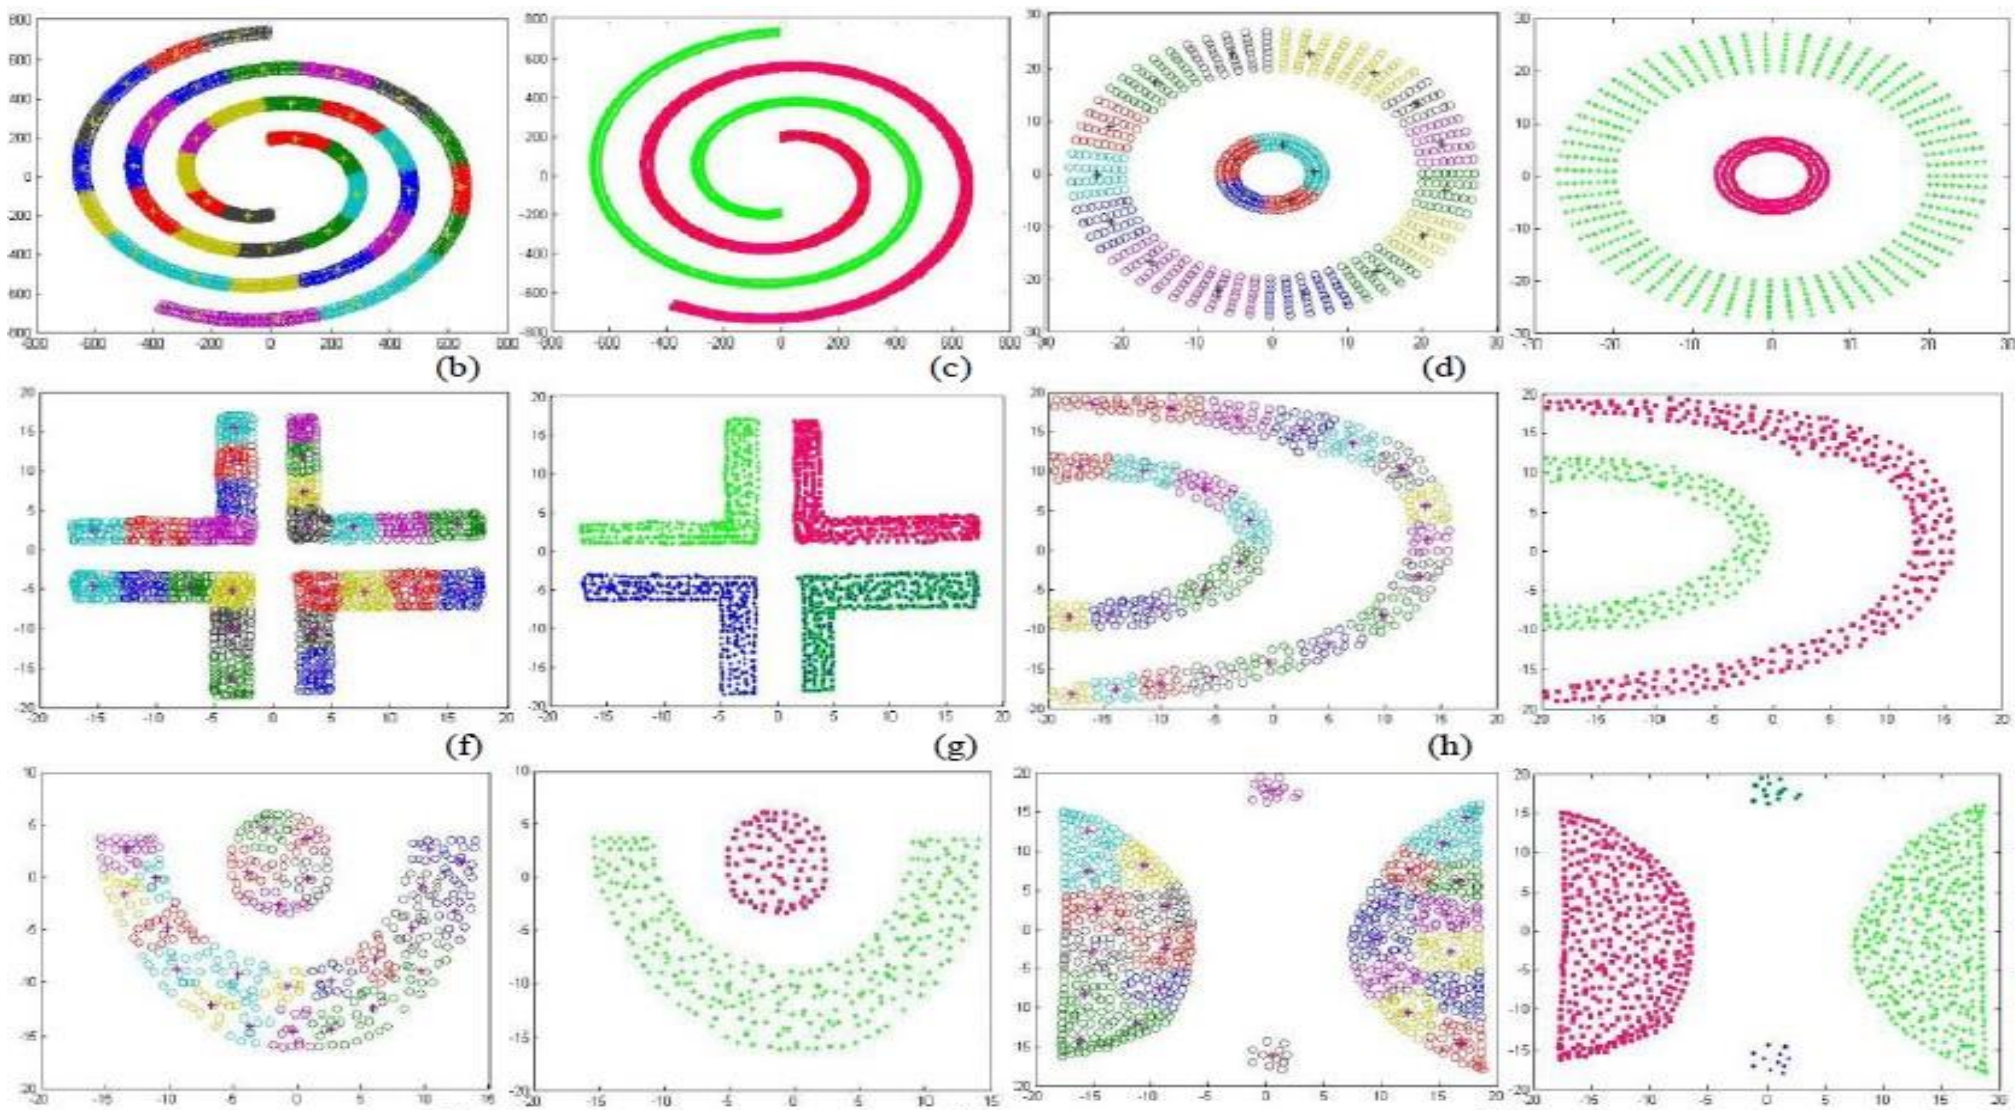
# Clustering: an easy task?



# Clustering: Challenging Clusters to discover!



# Clustering: Challenging Clusters to discover!





# Clustering: Definition & Formulation

- Clustering is the task of partitioning the data points into natural groups called clusters, such that points within a group are very similar, whereas points between different groups are as dissimilar as possible.
- $\mathcal{C} = \{C_1, C_2, \dots, C_k\}$
- $C_i = \{x_j | x_j \in C_i\}$
- $C_i \cap C_j = \emptyset \rightarrow$  **disjoint cluster no overlapping**
- Clustering is an unsupervised learning approach since it does not require a separate training dataset to learn the model parameters.

# Clustering: Different Data Types

- Numerical data
- Categorical data (including binary data)
  - *Discrete data, no natural order (e.g., gender, zip-code, and market-basket)*
- Text data: Popular in social media, Web, and social networks
  - *Features: High-dimensional, sparse, value corresponding to word frequencies*
- Multimedia data: Image, audio, video (e.g., on Flickr, YouTube)
  - *Multi-modal (often combined with text data)*
- Time-series data: Sensor data, stock markets, temporal tracking, forecasting, etc.
- Sequence data: Weblogs, biological sequences, system command sequences
- Stream data

# Clustering: Portioning Problem

- **Problem definition:** Given  $K$ , find a partition of  $K$  clusters that optimizes the chosen partitioning criterion
- A brute-force or exhaustive algorithm for finding a good clustering is simply to
  - *generate all possible partitions of  $n$  points into  $k$  clusters*
  - *evaluate clusters*
  - *Choose the best clusters*
- However, this is clearly infeasible, since there are  $O(k^n / k!)$  clusterings of  $n$  points into  $k$  groups.
- Global optimal: Needs to exhaustively enumerate all partitions
- Heuristic methods (i.e., greedy algorithms): K-Means, K-Medians, K-Medoids, etc.

# Clustering Paradigms

- Representative-based
  - ***K-means Clustering***
  - *Expectation-Maximization (EM) Algorithms*
- Hierarchical
- Density-based
- Graph-based
- Spectral clustering

# Clustering: K-means Algorithm

## **K-means Clustering**

- One of the simplest and most popular unsupervised machine learning algorithms
- K-means is a greedy algorithm that minimizes the squared distance of points from their respective cluster means
- It performs hard clustering, that is, each point is assigned to only one cluster.
- We also show how kernel K-means can be used for nonlinear clusters



# Clustering: K-means Algorithm

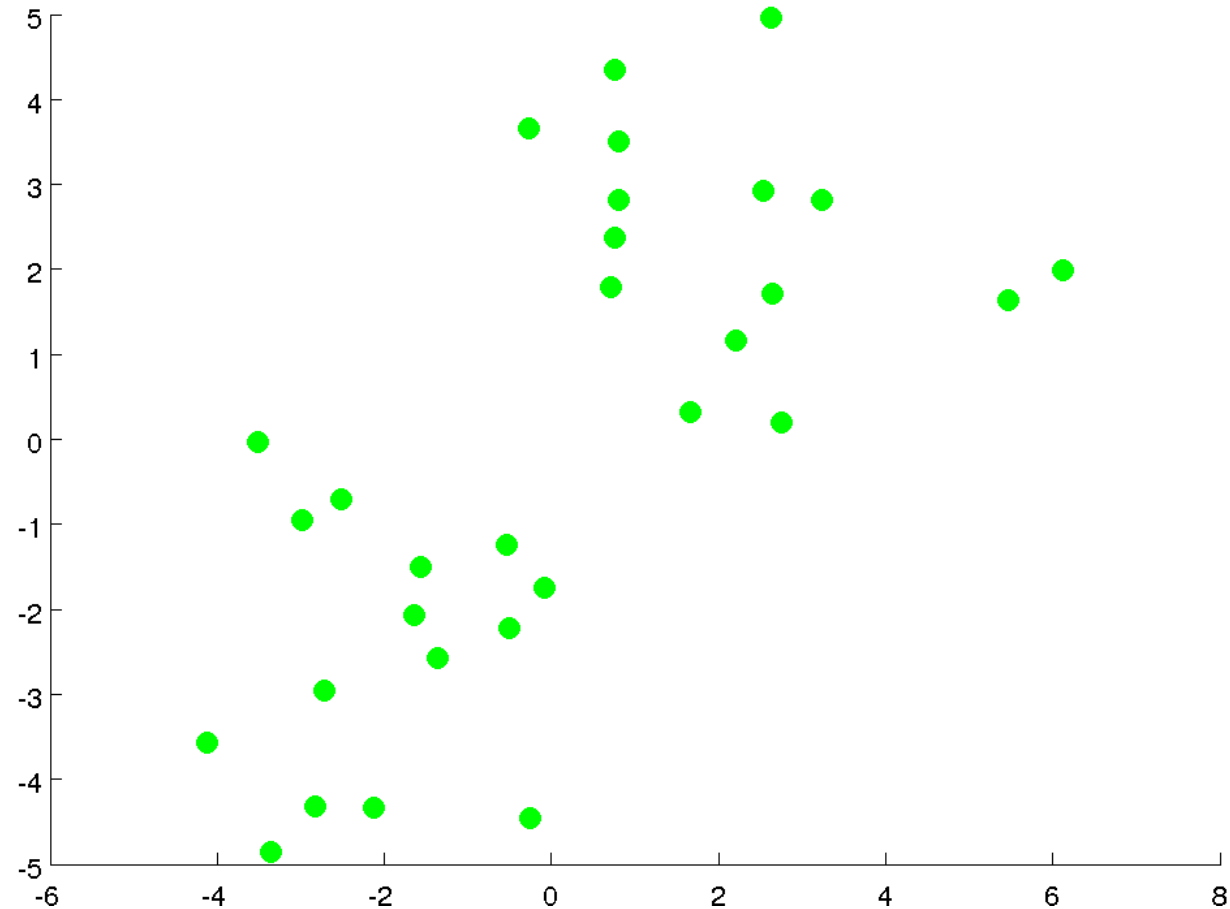
## K-means Clustering

### ➤ How does it work:

1. K is chosen as the number of clusters we wish to cluster our data into
2. Randomly choose centroid for each cluster
3. Iterate over the following two steps:
  - i. Cluster Assignment:  
Assign data to the cluster whose centroid is closest
  - ii. Centroid Adjustment:  
Move each centroid to the mean of data assigned to its cluster

# Clustering: K-means Algorithm

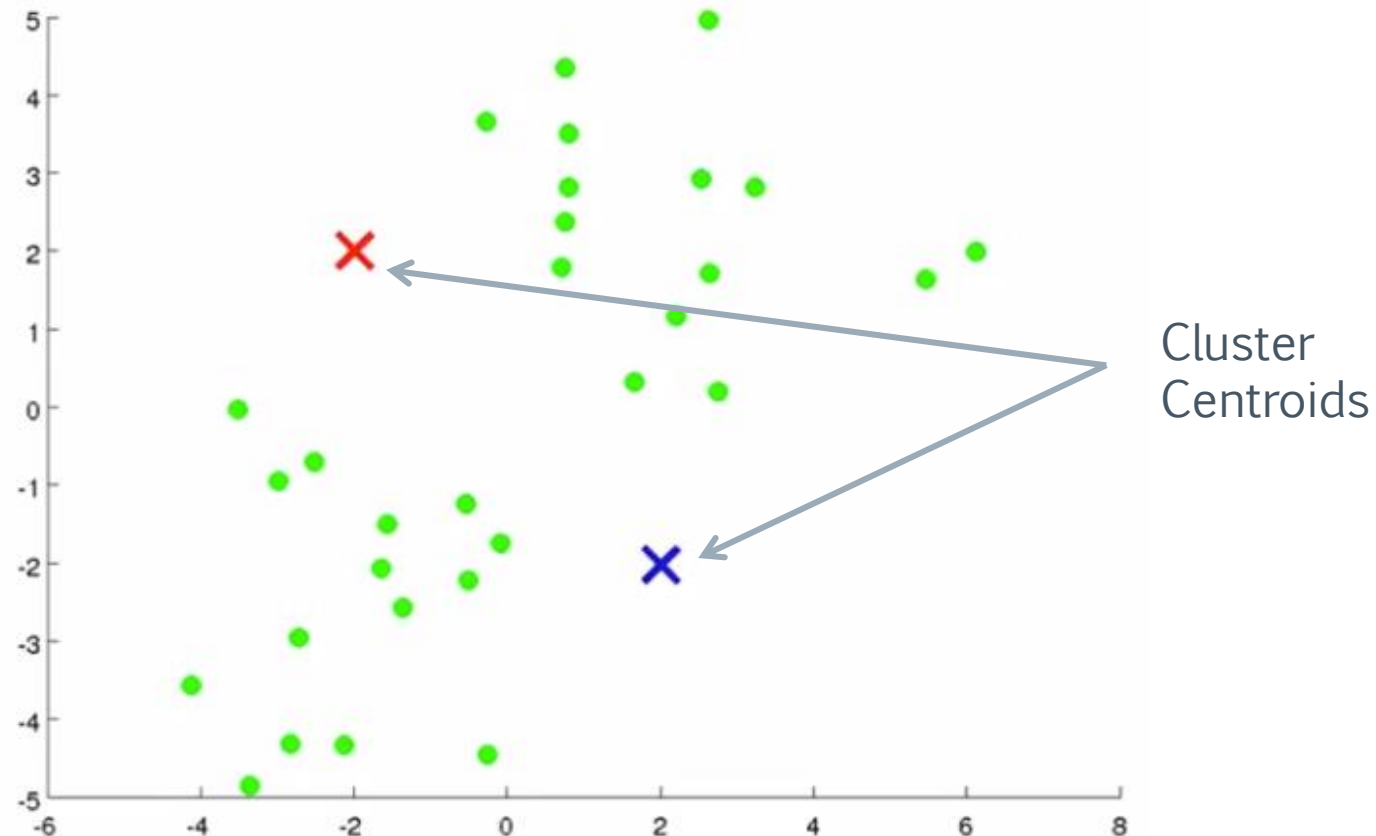
## K-means Clustering: Example (K=2)



# Clustering: K-means Algorithm

## K-means Clustering: Example (K=2)

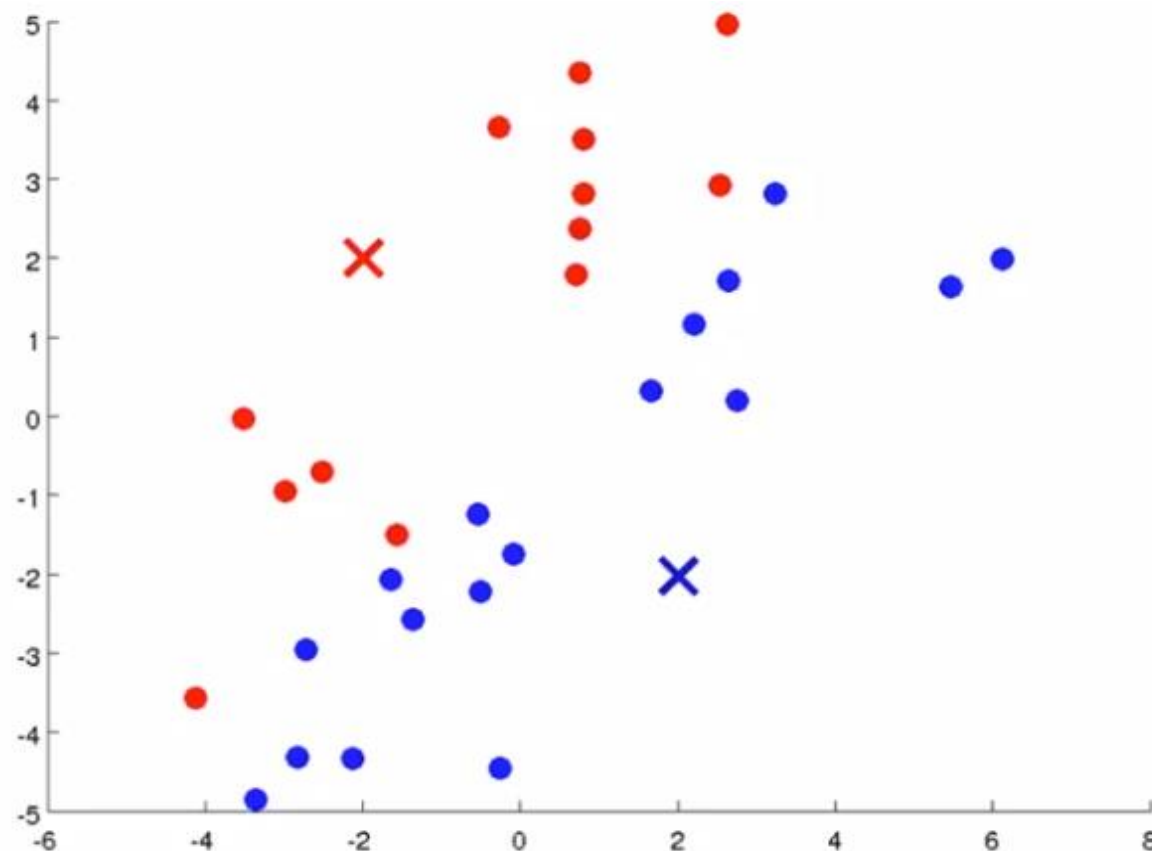
Step 1: Randomly assign cluster centroids



# Clustering: K-means Algorithm

## K-means Clustering: Example (K=2)

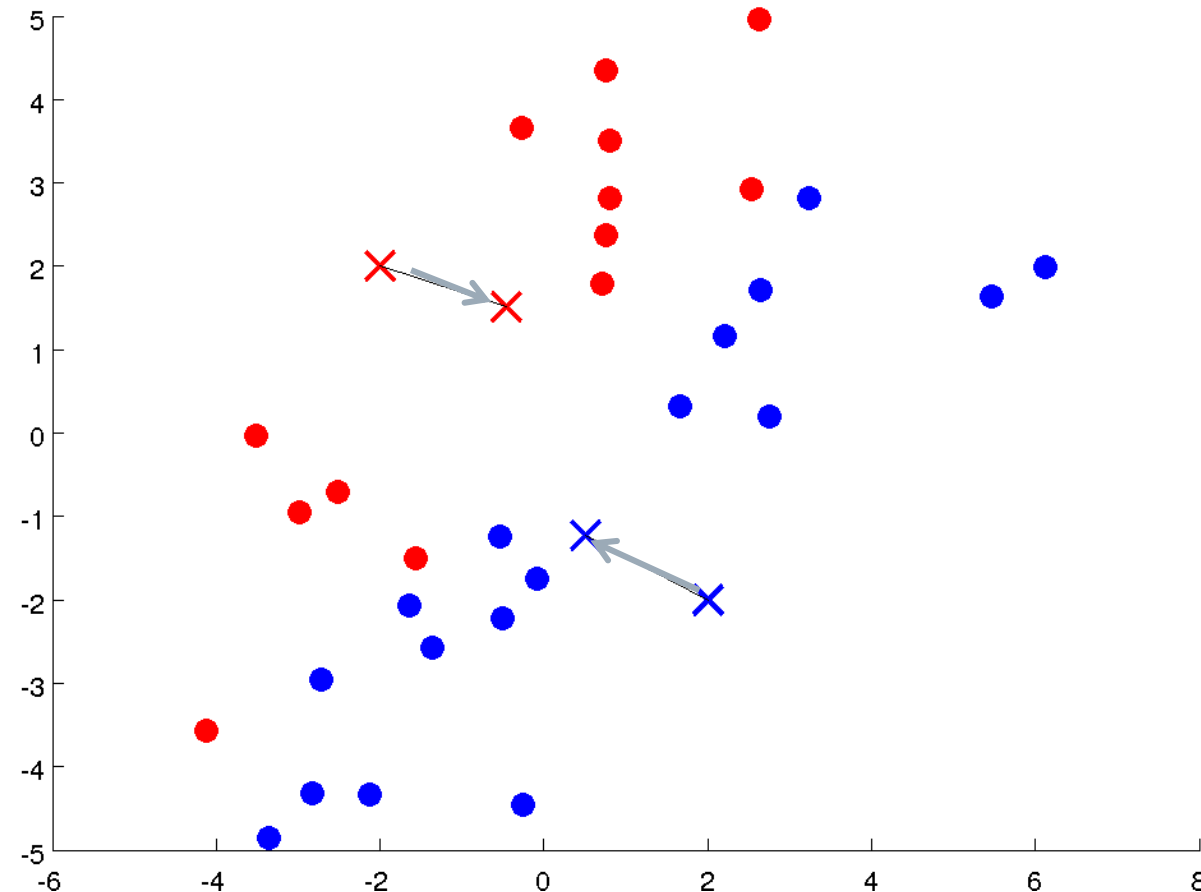
### Step 2.1: Cluster assignment



# Clustering: K-means Algorithm

## K-means Clustering: Example (K=2)

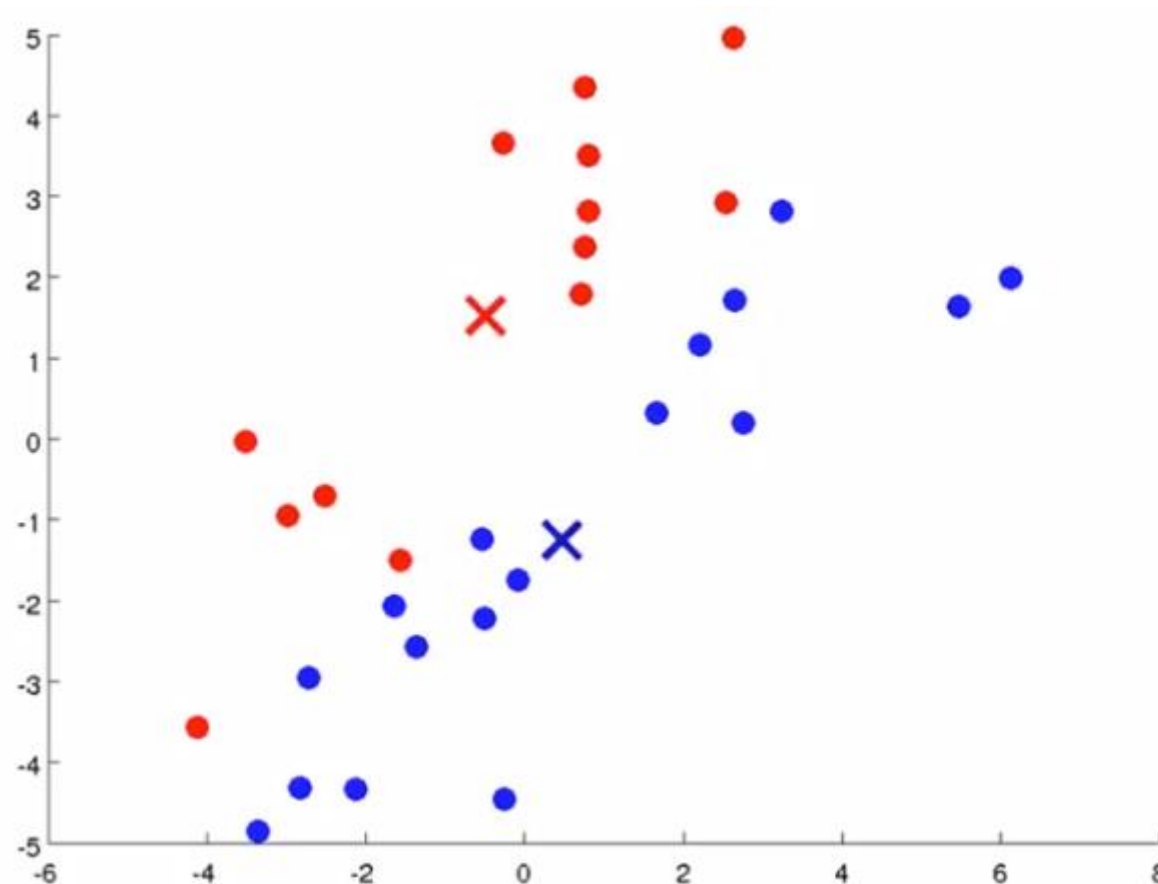
### Step 2.2: Centroid Adjustment



# Clustering: K-means Algorithm

## K-means Clustering: Example (K=2)

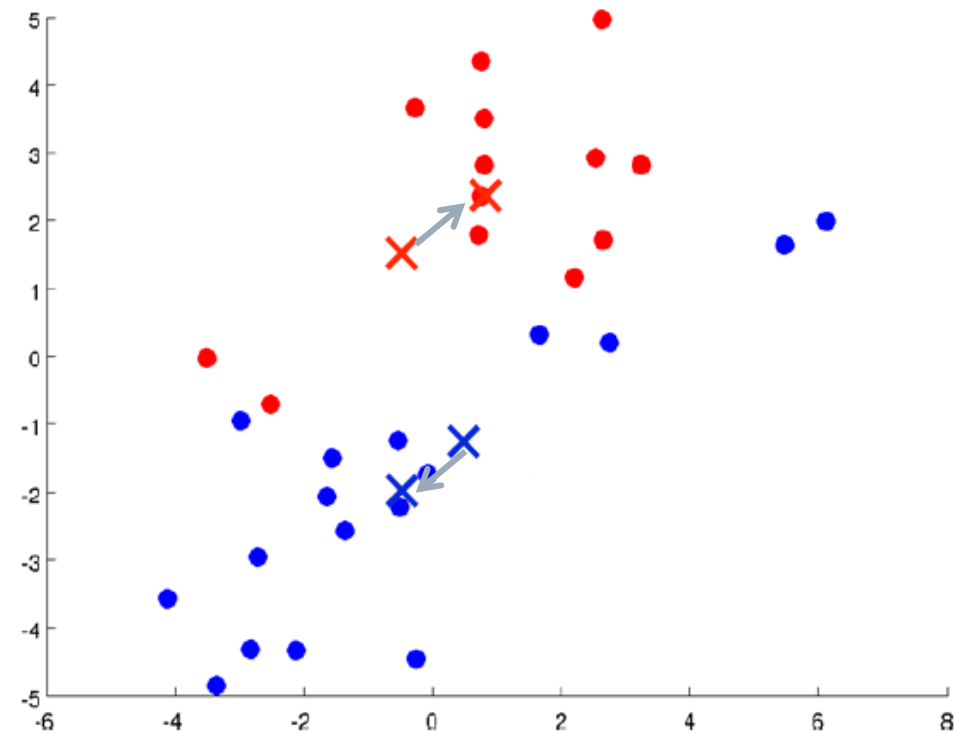
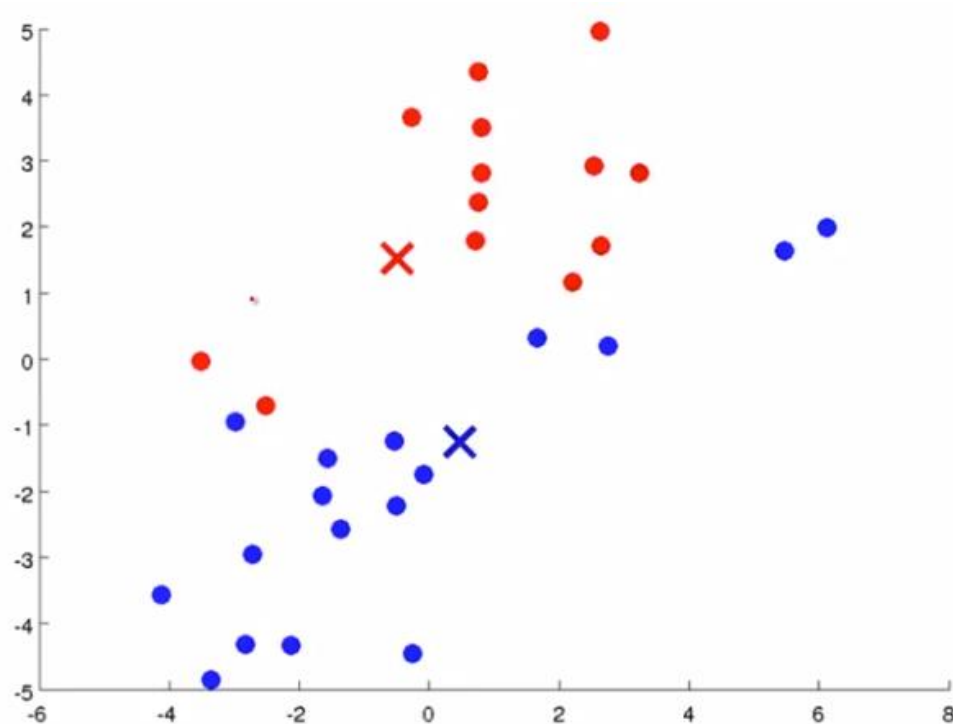
### Step 2.2: Centroid Adjustment



# Clustering: K-means Algorithm

## K-means Clustering: Example (K=2)

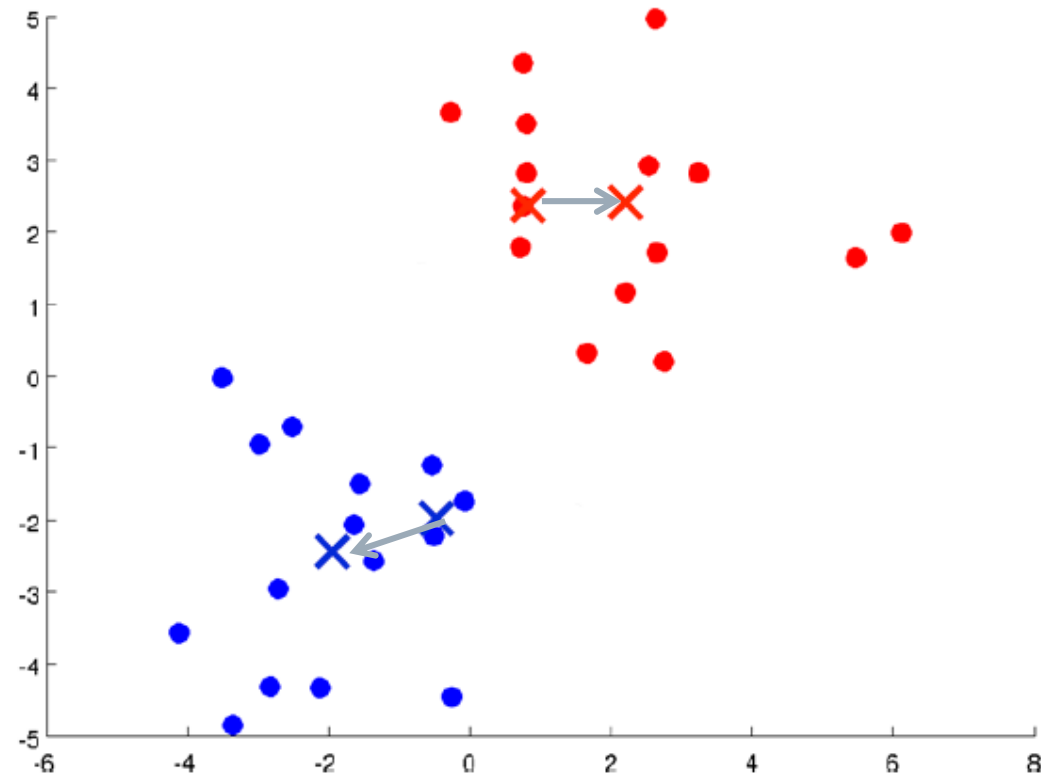
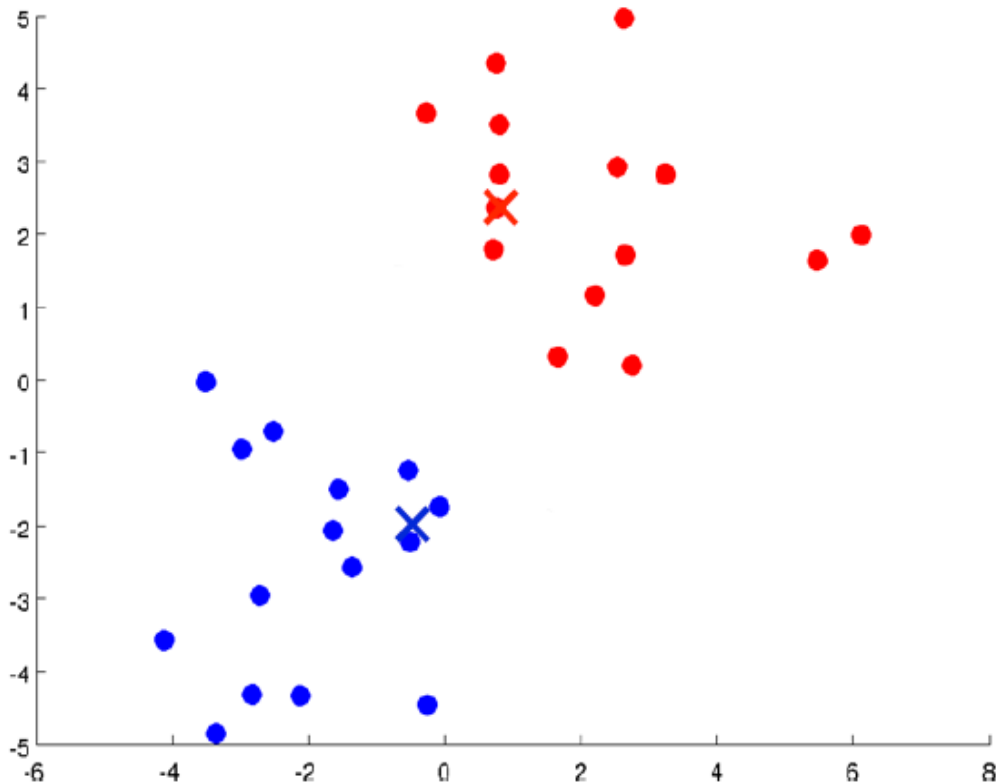
Repeat: Cluster Assignment & Centroid Adjustment



# Clustering: K-means Algorithm

## K-means Clustering: Example (K=2)

Repeat: Cluster Assignment & Centroid Adjustment

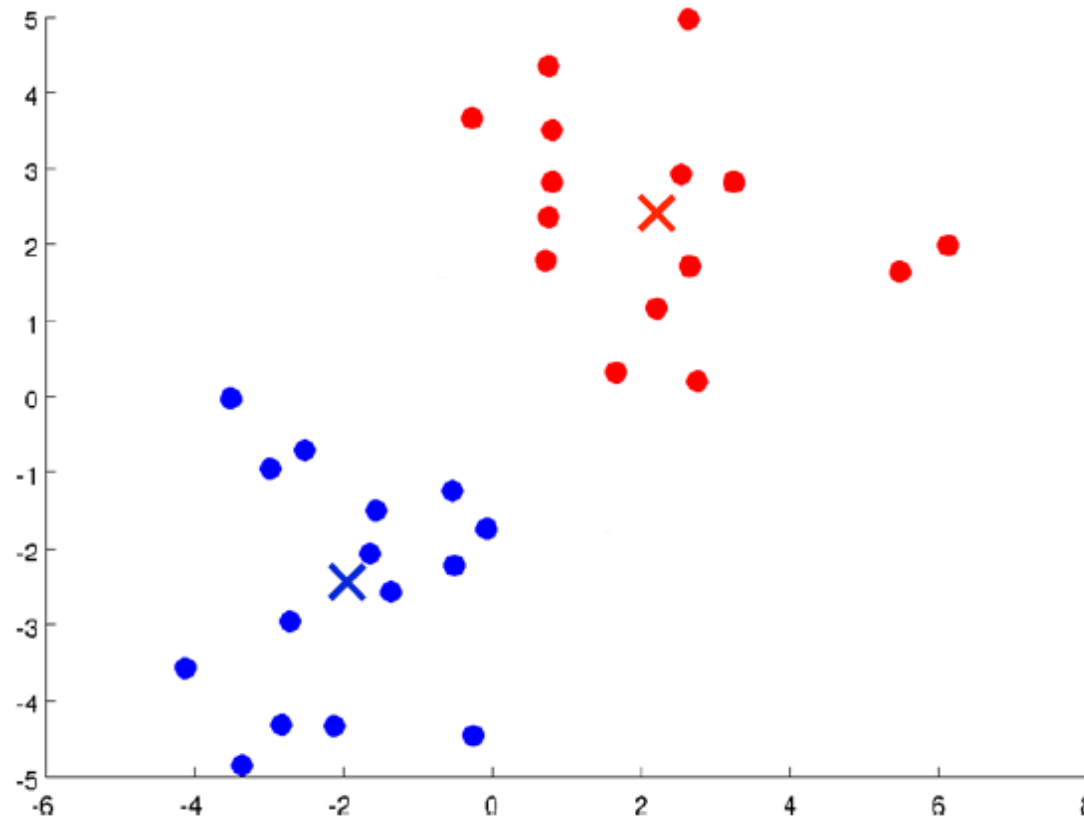




# Clustering: K-means Algorithm

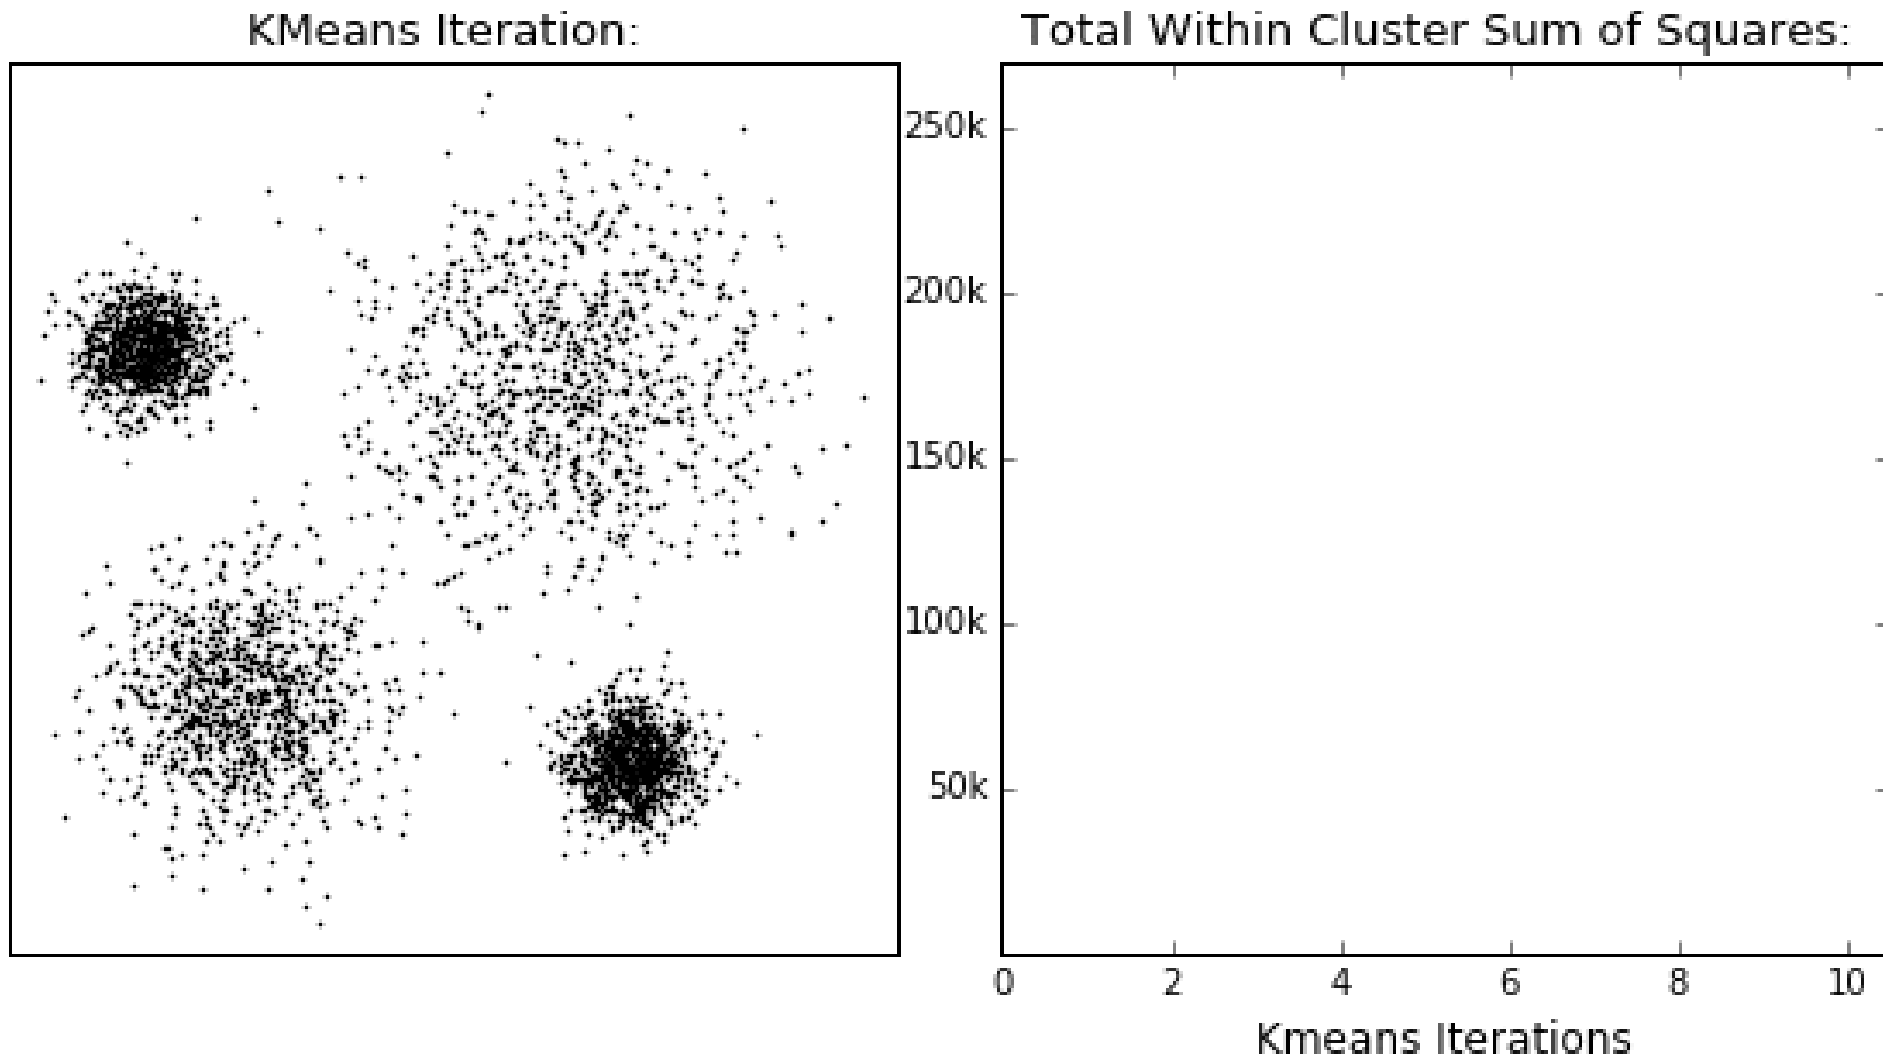
## K-means Clustering: Example (K=2)

Repeat: Cluster Assignment & Centroid Adjustment



Stop when  
there is no  
change in  
the cluster  
centroid

# Clustering: K-means Algorithm



# Clustering: K-means Algorithm

## Input:

- K (number of clusters)
- Training set  $\{x^{(1)}, x^{(2)}, x^{(3)}, \dots, x^{(m)}\}$

## Algorithm:

- Randomly initialize K cluster centroids  $\mu_1, \mu_2, \dots, \mu_K$
- Repeat until convergence
  - { for i=1 to m
    - $c^{(i)}$  = index (from 1 to K) of cluster centroid closest to  $x^{(i)}$   
Calculated as :  $\min_k \|x^{(i)} - \mu_k\|^2$
    - for k=1 to K
      - $\mu_k$  = average of points assigned to cluster k

# Clustering: K-means Algorithm

## Optimization Objective

$c^{(i)}$  = index of cluster (1,2,...,K) to which  $x^{(i)}$  is currently assigned

$\mu_K$  = cluster centroid k

$\mu_{c^{(i)}}$  = cluster centroid of cluster to which the example  $x^{(i)}$  has been assigned

$$\min_{\substack{c^{(1)}, \dots, c^{(m)} \\ \mu_1, \dots, \mu_K}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

where

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

and this objective function is also known as Distortion Function

# Example: K-means in 1D

- Consider the following one-dimensional data. Assume that we want to cluster the data into  $k = 2$  groups.  $\{2, 3, 4, 10, 11, 12, 20, 25, 30\}$ .

1) Initial centroids  $\mu_1 = 2$  and  $\mu_2 = 4$ .

2) Loop until convergence

A. First iteration

a) Cluster assignment, assigning each point to the closest mean:

$$C1 = \{2, 3\} \quad C2 = \{4, 10, 11, 12, 20, 25, 30\}$$

b) Centroid update, update the means

$$\mu_1 = \frac{2+3}{2} = \frac{5}{2} = 2.5 \quad \mu_2 = \frac{4+10+11+12+20+25+30}{7} = \frac{112}{7} = 16$$

B. Second iteration

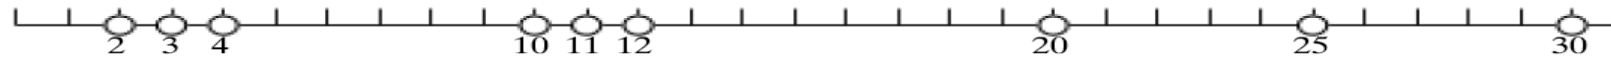
a) assigning each point to the closest mean:

$$C1 = \{2, 3, 4\} \quad C2 = \{10, 11, 12, 20, 25, 30\}$$

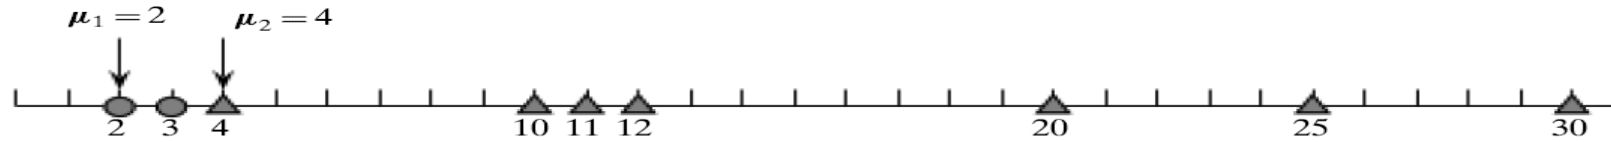
b) update the means

$$\mu_1 = \frac{2+3+4}{3} = \frac{9}{3} = 3 \quad \mu_2 = \frac{10+11+12+20+25+30}{6} = \frac{108}{6} = 18$$

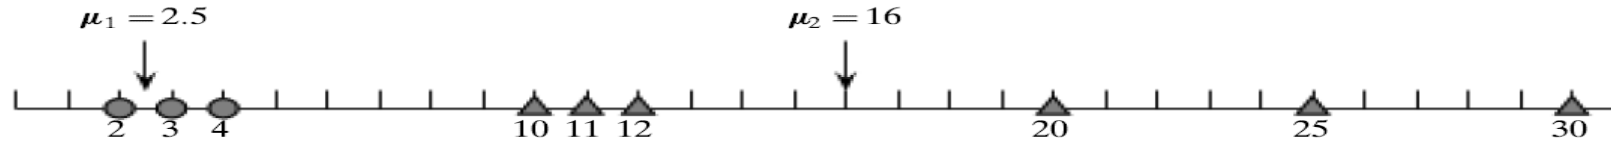
# Example: K-means in 1D



(a) Initial dataset



(b) Iteration:  $t = 1$



(c) Iteration:  $t = 2$



(e) Iteration:  $t = 4$



(f) Iteration:  $t = 5$  (converged)

$C_1 = \{2, 3, 4, 10, 11, 12\}$   $C_2 = \{20, 25, 30\}$  with centroids  $\mu_1 = 7$  and  $\mu_2 = 25$ .

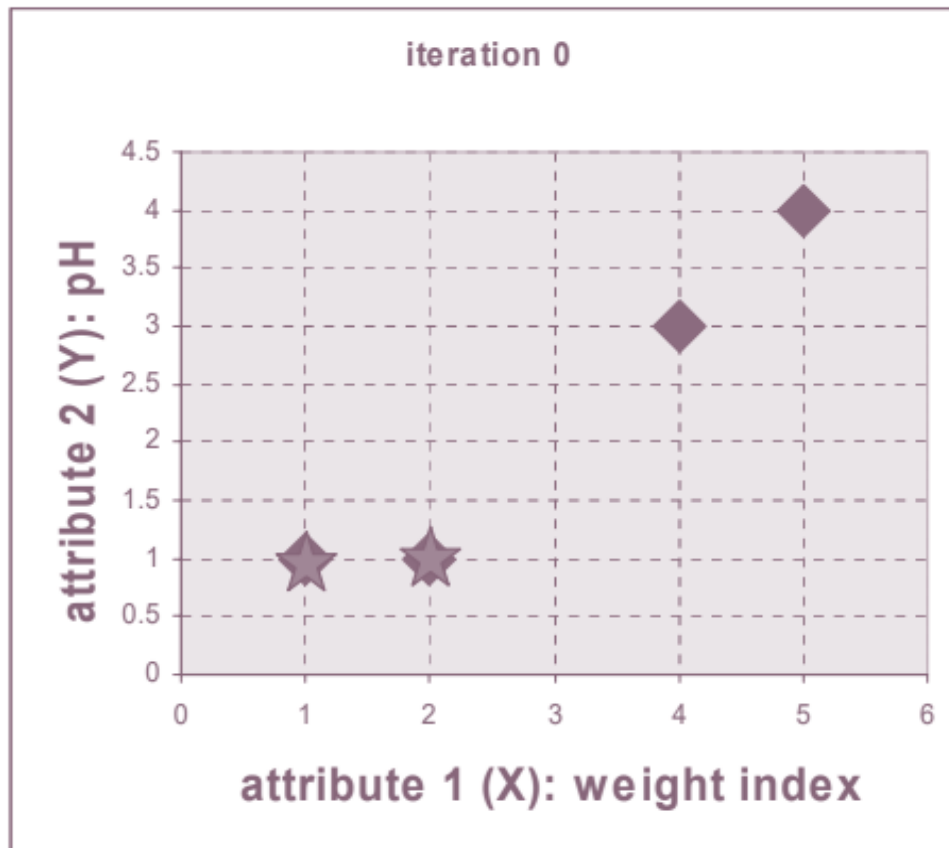
# Example: K-means in 2D

- Suppose we have several objects (4 types of medicines) and each object have two attributes or features as shown in table below. Our goal is to group these objects into  $K=2$  group of medicine based on the two features (pH and weight index).

Object	Attribute 1 (x1): weight index	Attribute 2 (x2): pH
A	1	1
B	2	1
C	4	3
D	5	4

# Example: K-means in 2D

- Each medicine represents one point with two features (X, Y) that we can represent it as coordinate in a feature space as shown in the figure.



Object	Attribute 1 (x1): weight index	Attribute 2 (x2): pH
A	1	1
B	2	1
C	4	3
D	5	4



# Example: K-means in 2D

- Initial value of centroids:  $\mu_1 = (1,1)$  and  $\mu_2 = (2,1)$
- $\mu_1 = (1,1)$  with Medicine C
- $= \sqrt{(4 - 1)^2 + (3 - 1)^2}$
- $= \sqrt{(3)^2 + (2)^2}$
- $= \sqrt{9 + 4}$
- $\sqrt{13} = 3.61$

Object	Attribute 1 (x1): weight index	Attribute 2 (x2): pH
A	1	1
B	2	1
C	4	3
D	5	4

# Example: K-means in 2D

## ➤ Cluster assignment

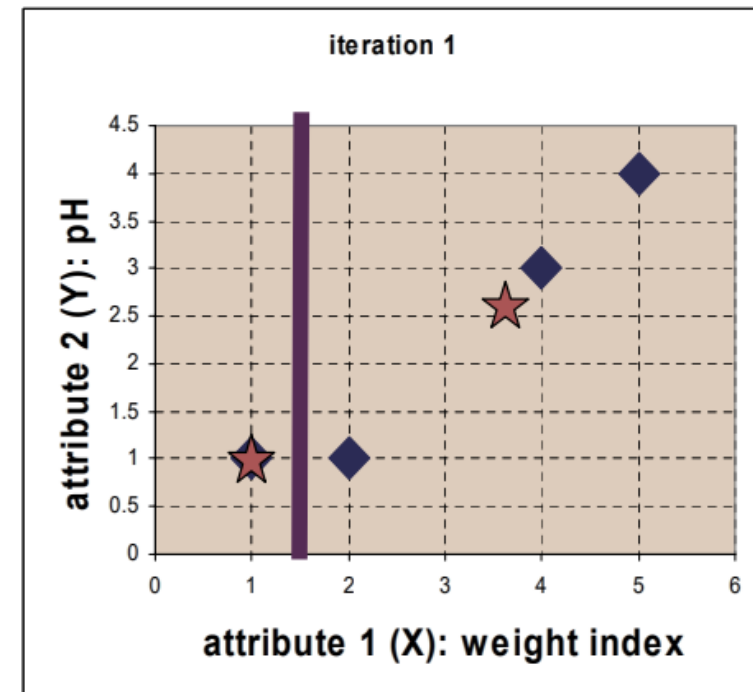
	$\mu_1$	$\mu_2$	Cluster
A	0	1	1
B	1	0	2
C	3.605551	2.828427	2
D	5	4.242641	2

## ➤ Update Centroid

➤  $\mu_1 = (1, 1)$

➤  $\mu_2 = \frac{2+4+5}{3}, \frac{1+3+4}{3} = \frac{11}{3}, \frac{8}{3}$

Object	weight index	pH
A	1	1
B	2	1
C	4	3
D	5	4



## Example: K-means in 2D (2)

- Use the k-means algorithm and Euclidean distance to cluster the following 8 examples into 3 clusters:
- $A1=(2,10)$ ,  $A2=(2,5)$ ,  $A3=(8,4)$ ,  $A4=(5,8)$ ,  $A5=(7,5)$ ,  $A6=(6,4)$ ,  $A7=(1,2)$ ,  $A8=(4,9)$ . The distance matrix based on the Euclidean distance is given below

	A1	A2	A3	A4	A5	A6	A7	A8
A1	0	$\sqrt{25}$	$\sqrt{36}$	$\sqrt{13}$	$\sqrt{50}$	$\sqrt{52}$	$\sqrt{65}$	$\sqrt{5}$
A2		0	$\sqrt{37}$	$\sqrt{18}$	$\sqrt{25}$	$\sqrt{17}$	$\sqrt{10}$	$\sqrt{20}$
A3			0	$\sqrt{25}$	$\sqrt{2}$	$\sqrt{2}$	$\sqrt{53}$	$\sqrt{41}$
A4				0	$\sqrt{13}$	$\sqrt{17}$	$\sqrt{52}$	$\sqrt{2}$
A5					0	$\sqrt{2}$	$\sqrt{45}$	$\sqrt{25}$
A6						0	$\sqrt{29}$	$\sqrt{29}$
A7							0	$\sqrt{58}$
A8								0

# Application case: K-Means for Segmentation

Goal of Segmentation is to partition an image into regions each of which has reasonably homogenous visual appearance.

**Original**



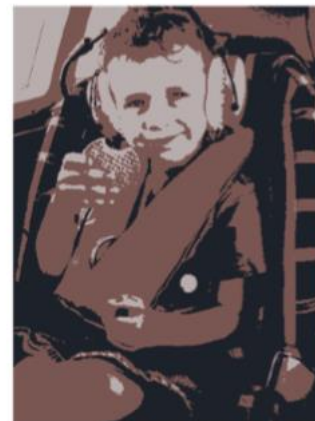
**K=2**



**K=3**



**K=10**

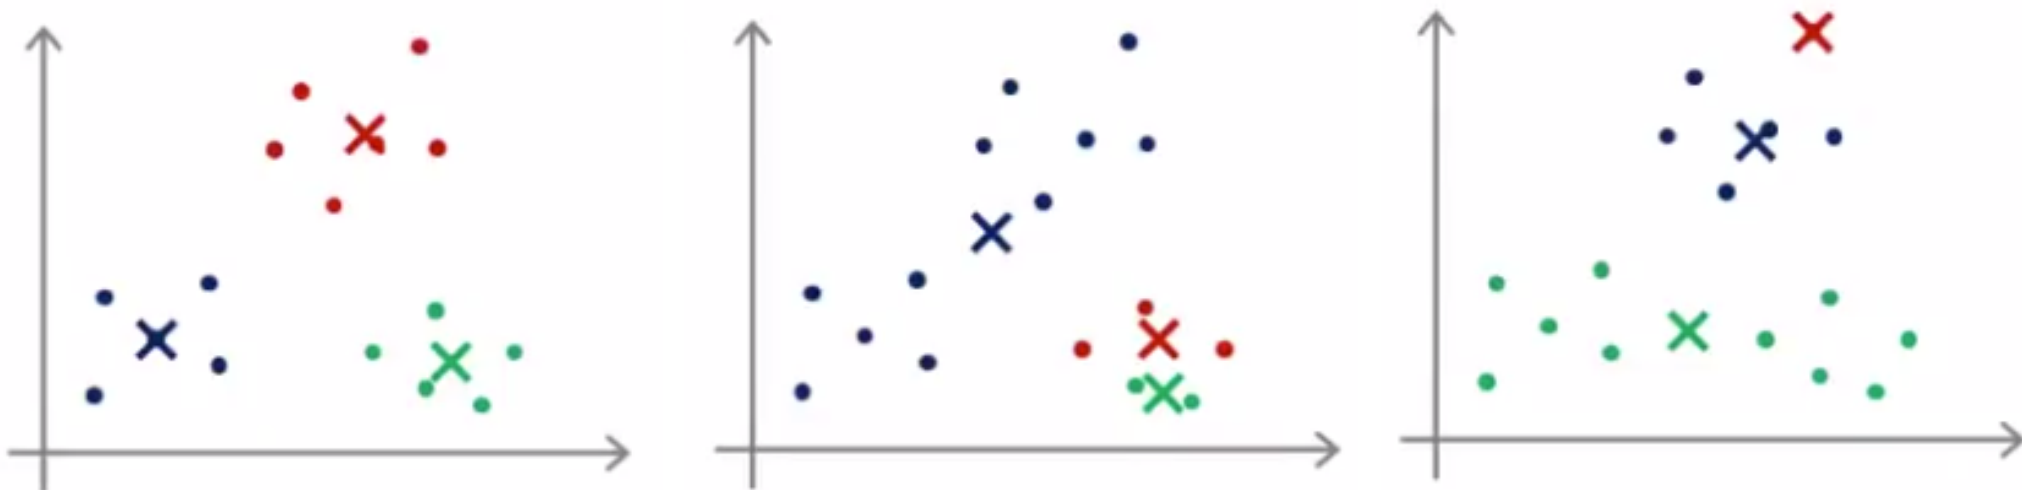


# K-means Algorithm: Random Initialization

How to randomly initialize cluster centroids?

1. Choose  $K < m$
2. Randomly pick  $K$  training examples
3. Set  $\mu_1, \dots, \mu_K$  equal to these  $K$  examples

Seems correct, but will it always work?



**No**, because K-means can get stuck at different local optimas

# K-means Algorithm: Random Initialization

## **Solution:**

Instead of initializing K-means once and hoping that it works;

- Initialize and run K-means many times, and use the solution that gives best local or global optima as possible.

## **So we do the following:**

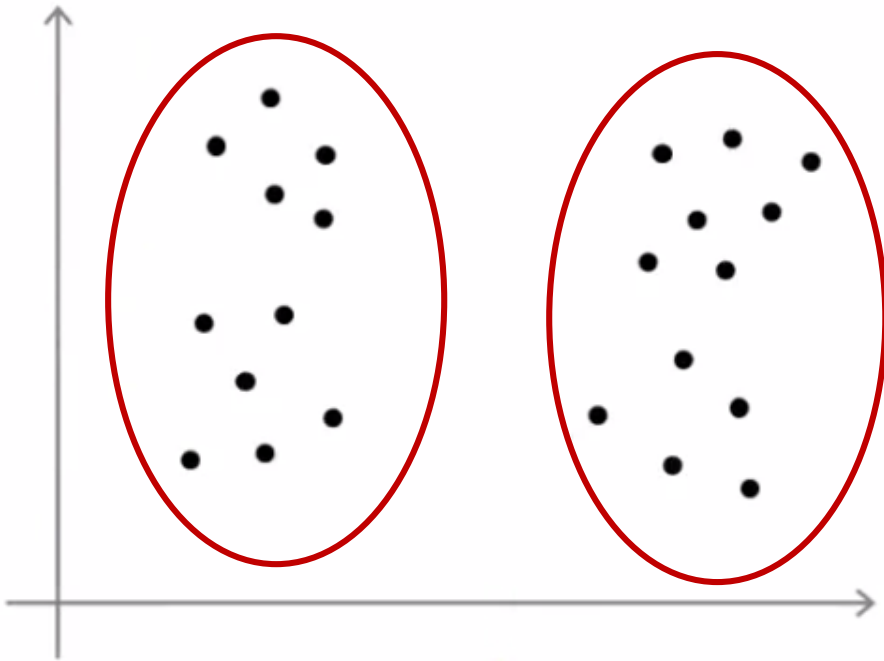
for i=1 to 100

```
{  
    Randomly initialize K-means  
    Run K-means to get  $c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K$   
    Compute cost function (distortion)  $J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$   
}
```

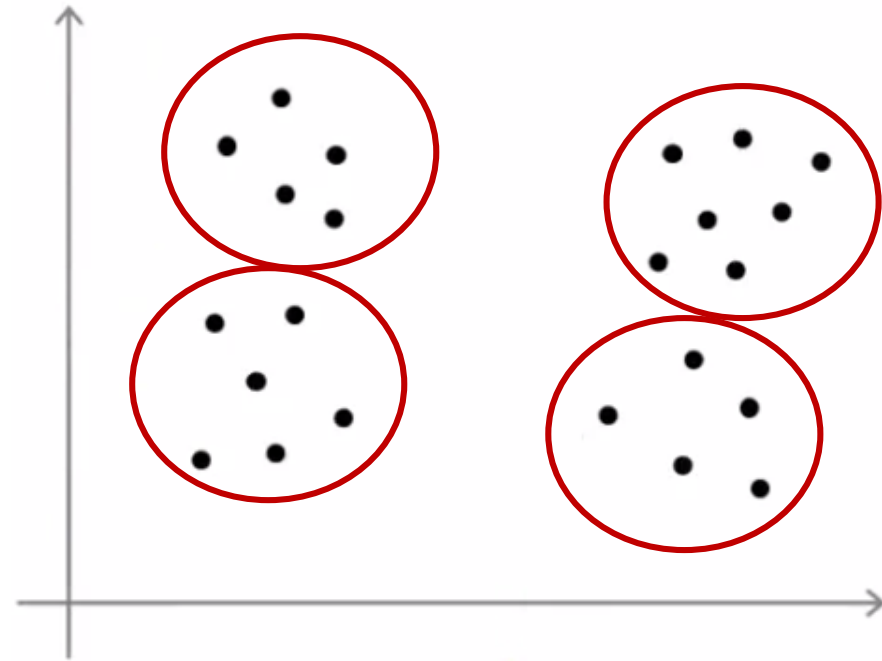
- **Pick clustering that gave lowest cost  $J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$**

# K-means Algorithm: Choosing number of clusters

K=2



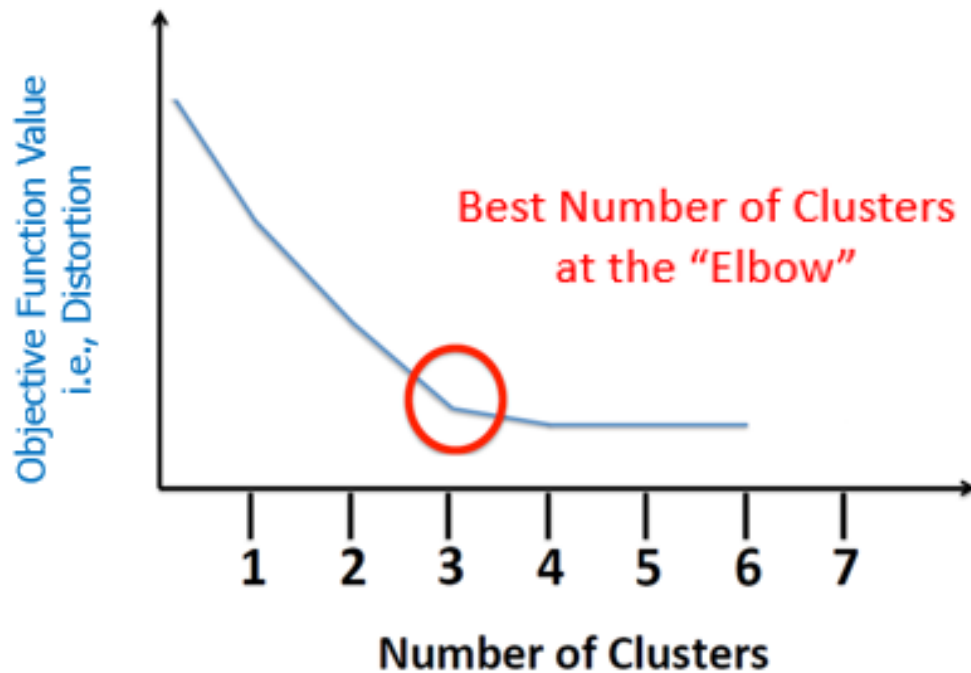
K=4



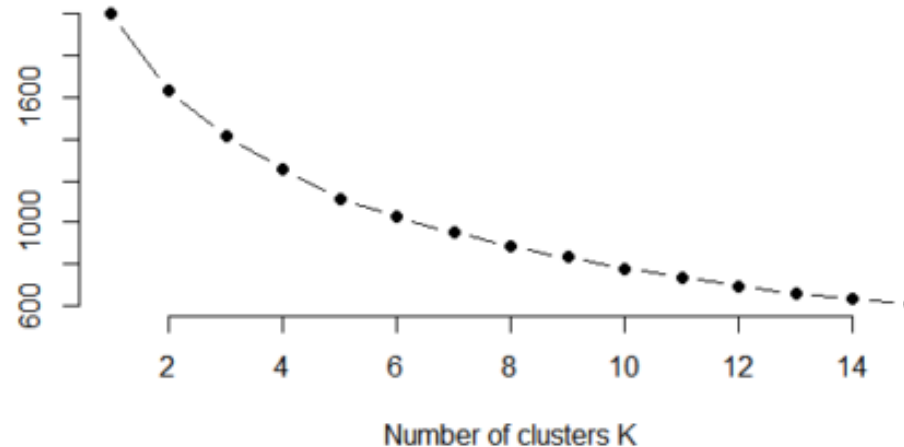
# K-means Algorithm: Choosing number of clusters

One possible way is using **Elbow Method**

- Plotting cost function  $J$  verses number of clusters



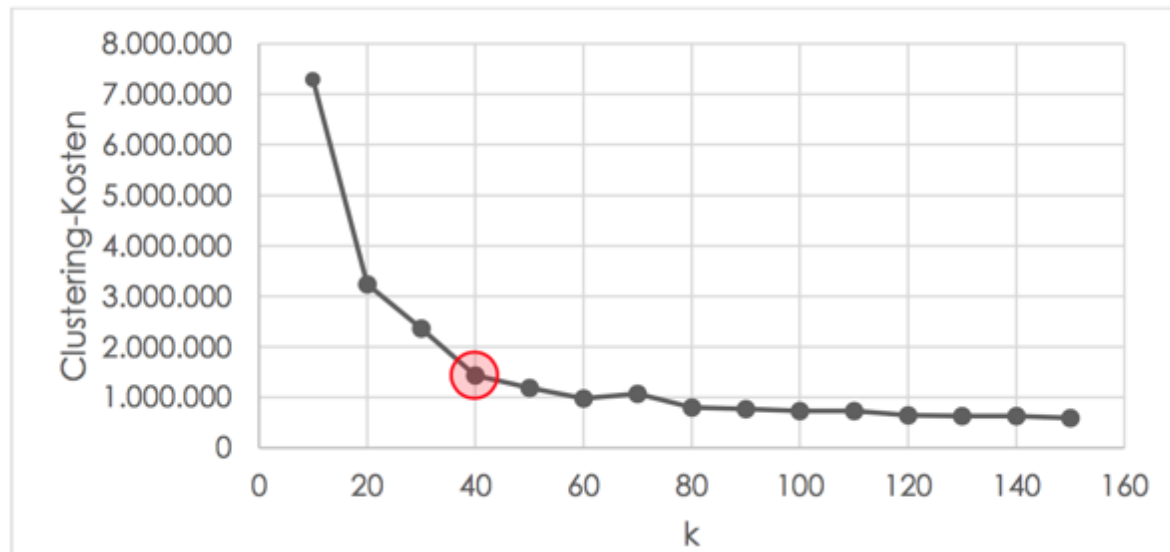
But, you may not always find an "elbow"...





# K-means Algorithm: Choosing number of clusters

- Using the elbow method we run k-means clustering for a range of values of  $k$ . (e.g. 1 to 150).
- For each value of  $k$  we then compute the sum of squared errors (SSE) and add both into a line plot.
- Illustration 1 shows an exemplary curve of a range of values of  $k$  and the corresponding SSE.

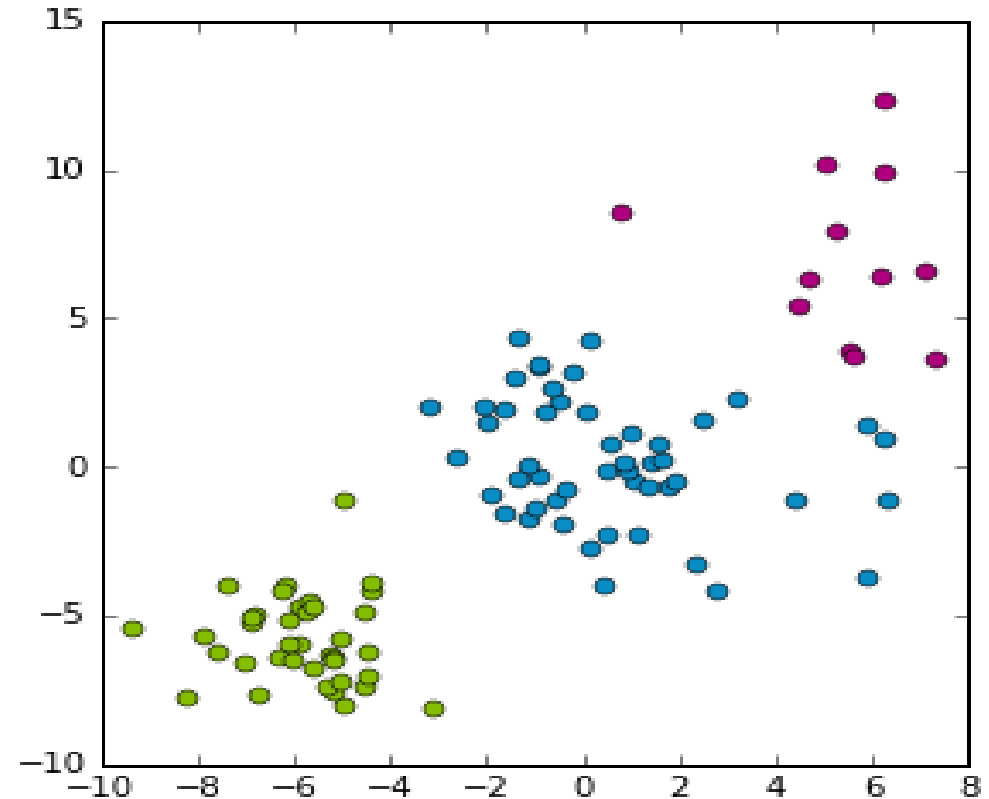
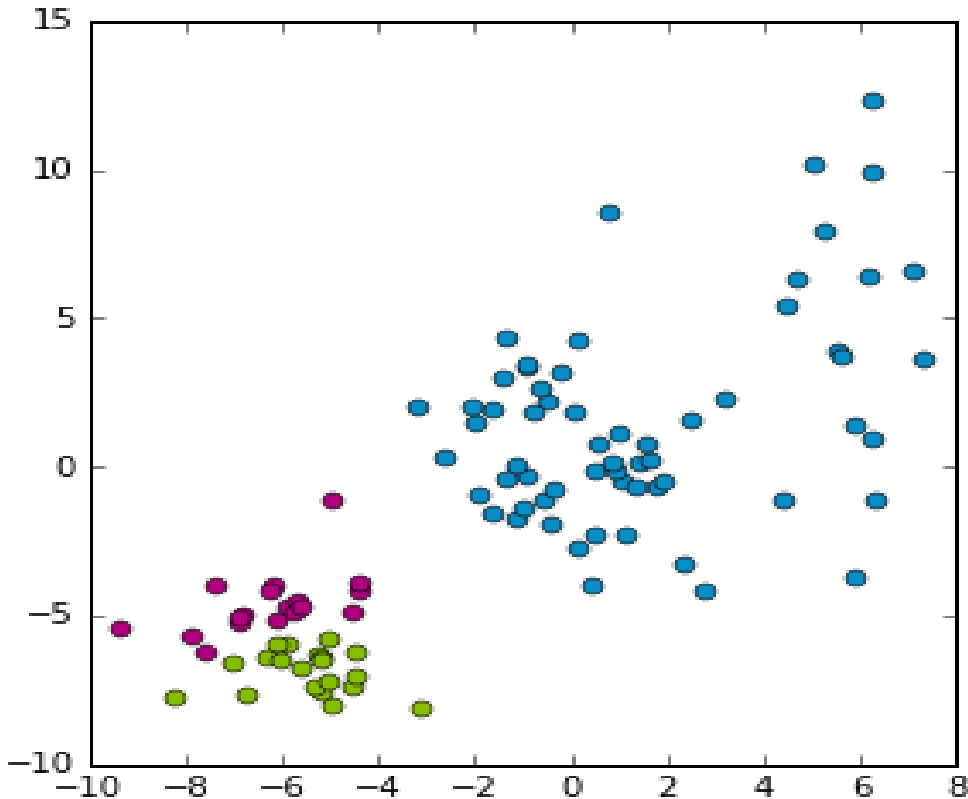


# K-means Algorithm: Evaluation

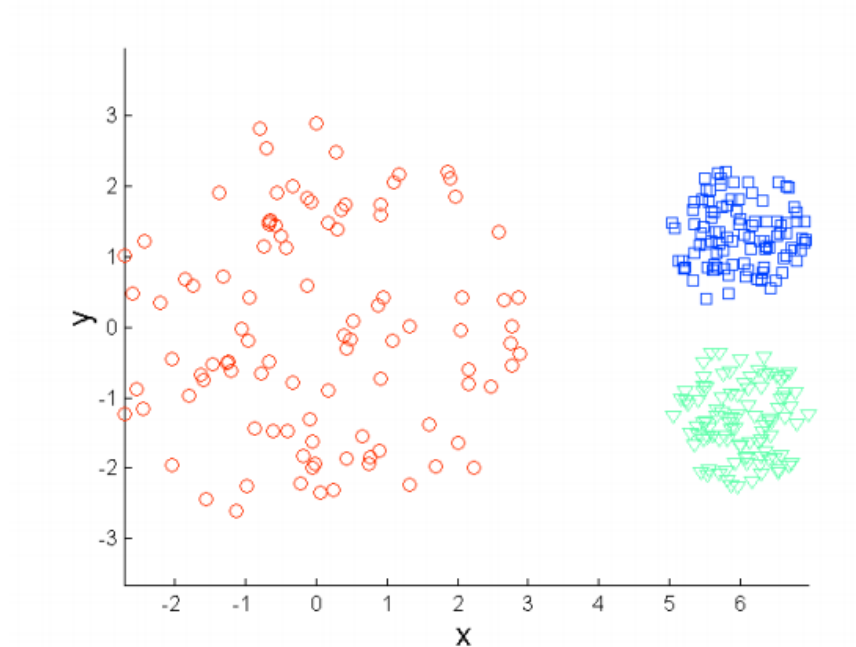
- Guaranteed to converge in a finite number of iterations
- Running time per iteration:
  1. Assign data points to closest cluster center:  
 $O(KN)$  time
  2. Change the cluster center to the average of its assigned points  $O(N)$

# K-means Algorithm: Assessing Clustering Quality

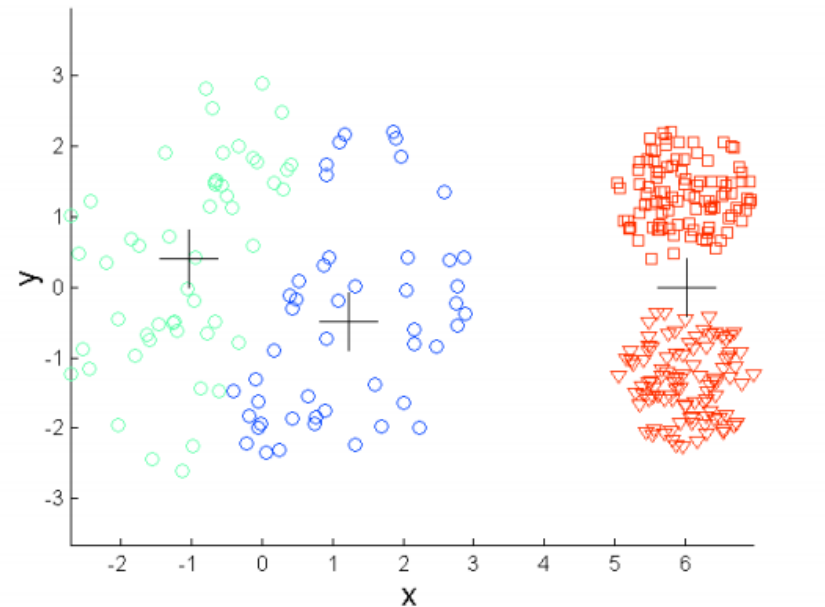
➤ Which clustering is better?



# K-means Algorithm: Assessing Clustering Quality



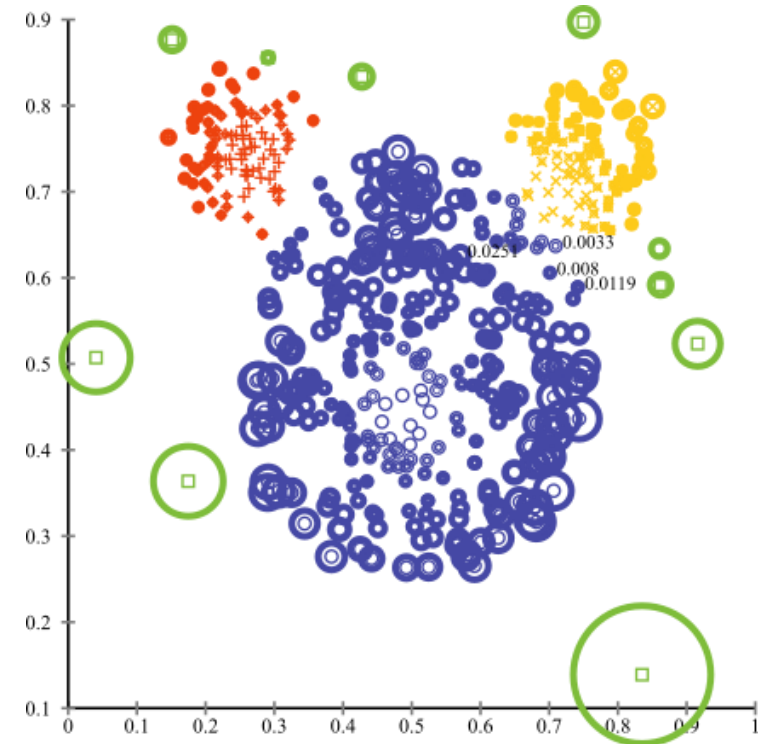
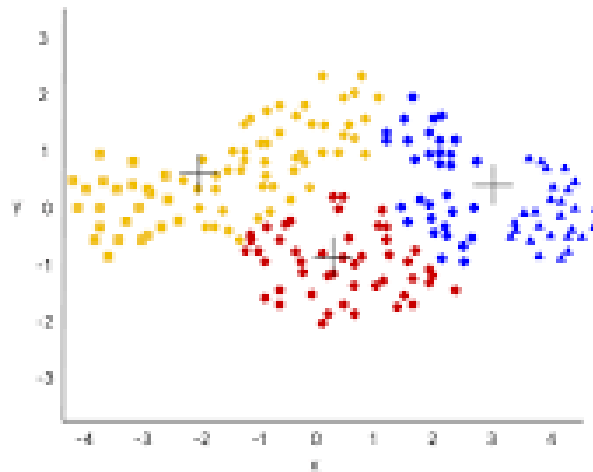
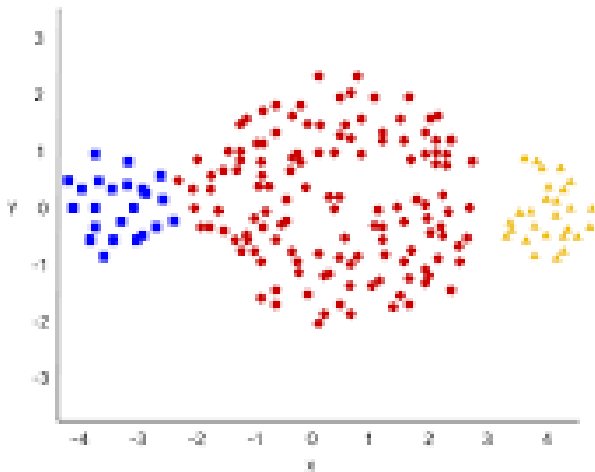
Original Points



K-means (3 Clusters)

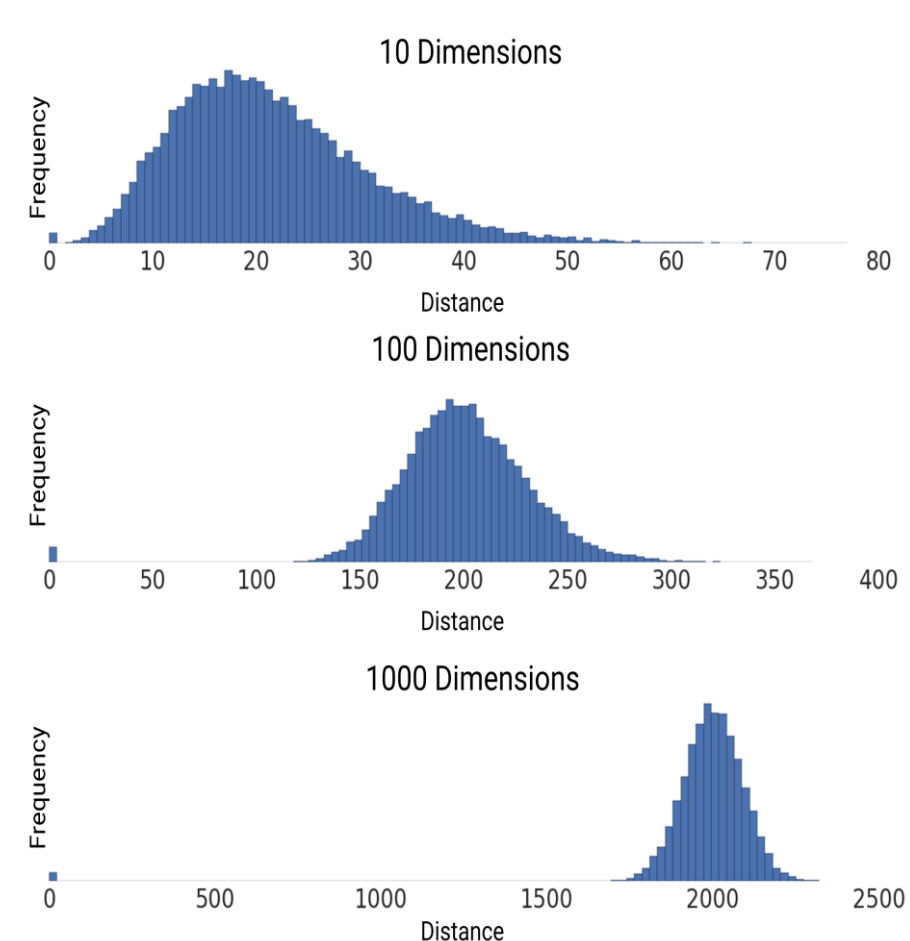
# K-means Algorithm: Clustering with Outliers

- Centroids can be dragged by outliers, or outliers might get their own cluster instead of being ignored.
- Consider removing or clipping outliers before clustering.



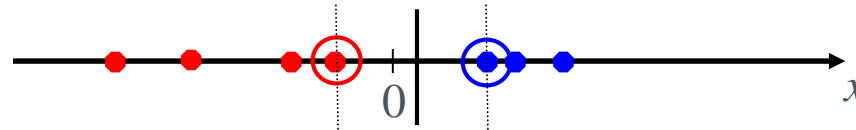
# K-means Algorithm: Curse of Dimensionality

- These plots show how the ratio of the standard deviation to the mean of distance between examples decreases as the number of dimensions increases.
- This convergence means k-means becomes less effective at distinguishing between examples.
- This negative consequence of high-dimensional data is called the curse of dimensionality.

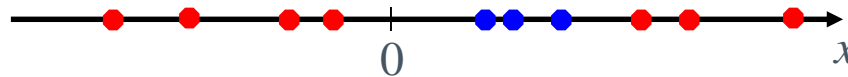


# K-means Algorithm: Non-linear Separators

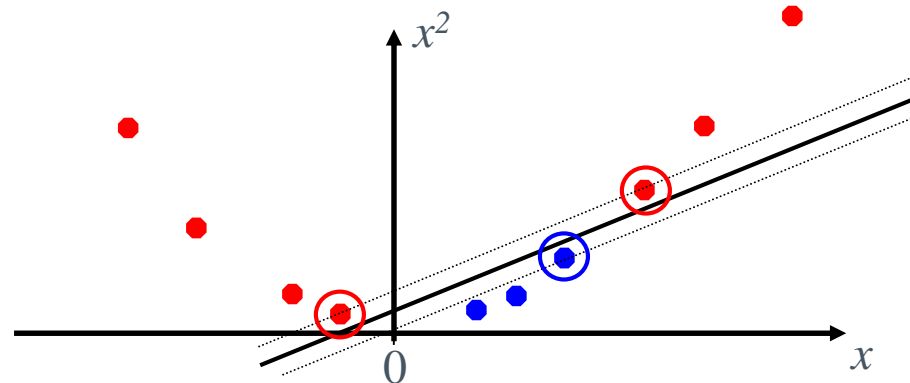
- Data that is linearly separable works out great for linear decision rules:



- But what are we going to do if the dataset is just too hard?

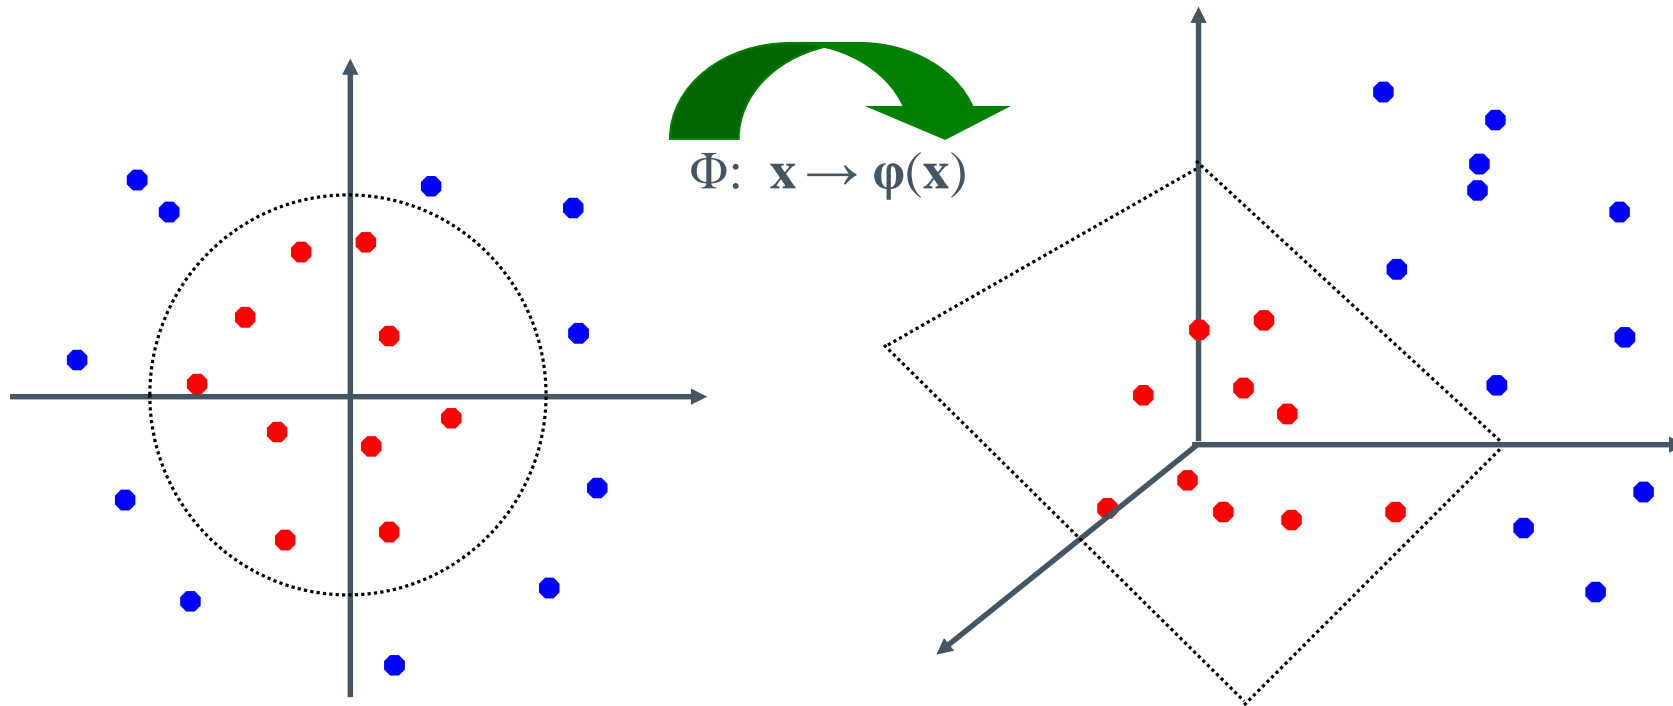


- How about... mapping data to a higher-dimensional space:



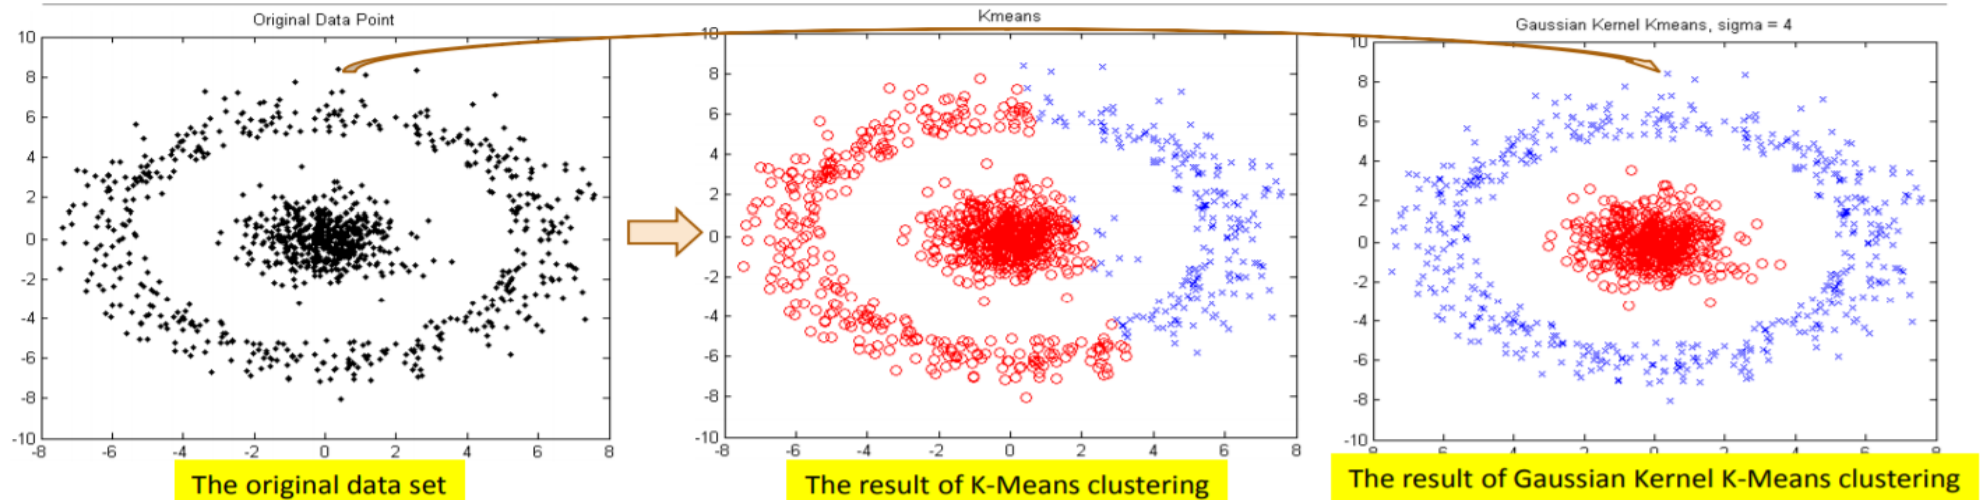
# K-means Algorithm: Non-linear Separators

- General idea: the original feature space can always be mapped to some higher-dimensional feature space where the training set is separable





# K-means clustering with Kernels



- ❑ The above data set cannot generate quality clusters by K-Means since it contains non-convex clusters
- ❑ Gaussian RBF Kernel transformation maps data to a kernel matrix  $K$  for any two points  $x_i, x_j$ :  $K_{x_i x_j} = \phi(x_i) \bullet \phi(x_j)$  and Gaussian kernel:  $K(\mathbf{x}_i, \mathbf{x}_j) = e^{-\|\mathbf{x}_i - \mathbf{x}_j\|^2 / 2\sigma^2}$
- ❑ K-Means clustering is conducted on the mapped data, generating quality clusters