

# Machine Learning Fundamentals – DTSC102

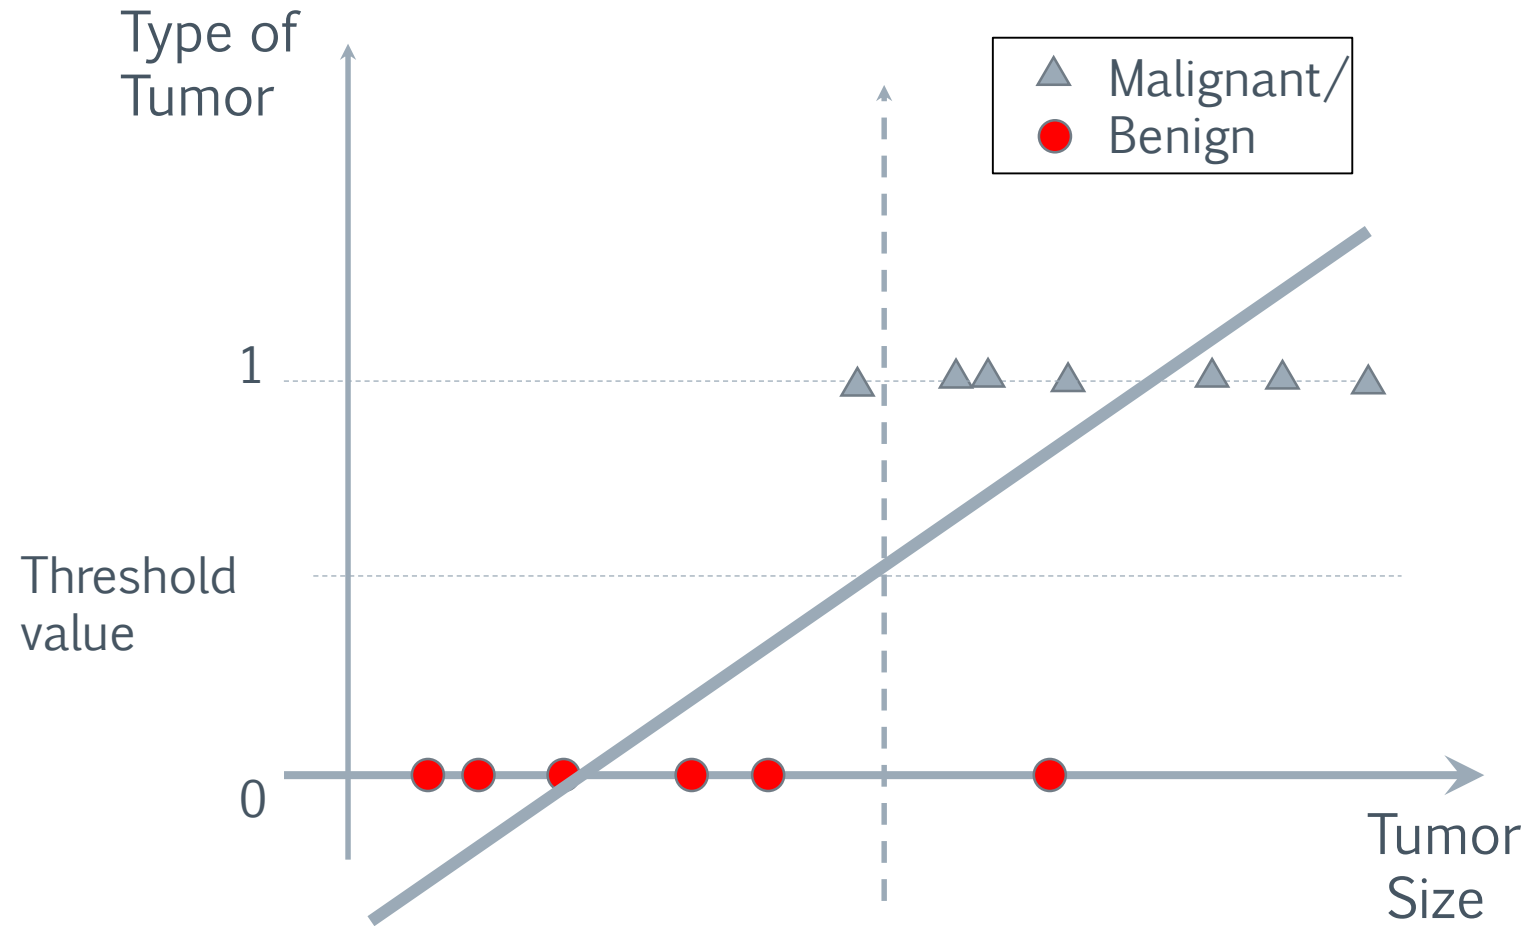
## Lecture 6 Logistic Regression

Course Instructor: Dr.-Ing. Maggie Mashaly  
maggie.ezzat@guc.edu.eg  
C3.220

# Contents

- Classification
- Logistic Regression
- Decision Boundaries
- Cost Function

# Classification

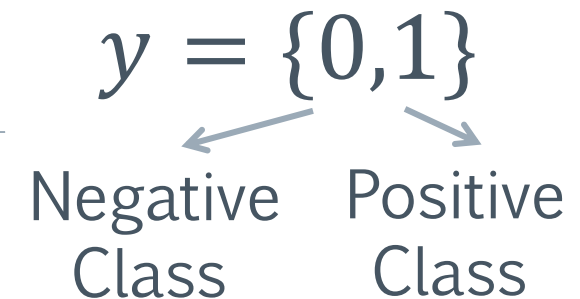


# Classification

- One form of Supervised Learning
- Goal: Learn a mapping from inputs  $x$  to outputs  $y$ , where  $y=\{1,2,...,C\}$ , with  $C$  being the number of classes

- Examples:

- Emails: Spam/Not Spam
- Online Transactions: Fraudulent (Yes/No)
- Tumor: Malignant/Benign



- Approach: Predict a hypothesis function

$$0 \leq h_{\theta}(x) \leq 1$$

# Logistic Regression

Approach: Predict a hypothesis function

$$0 \leq h_{\theta}(x) \leq 1$$

Solution:

- Use Logistic/Sigmoid Function

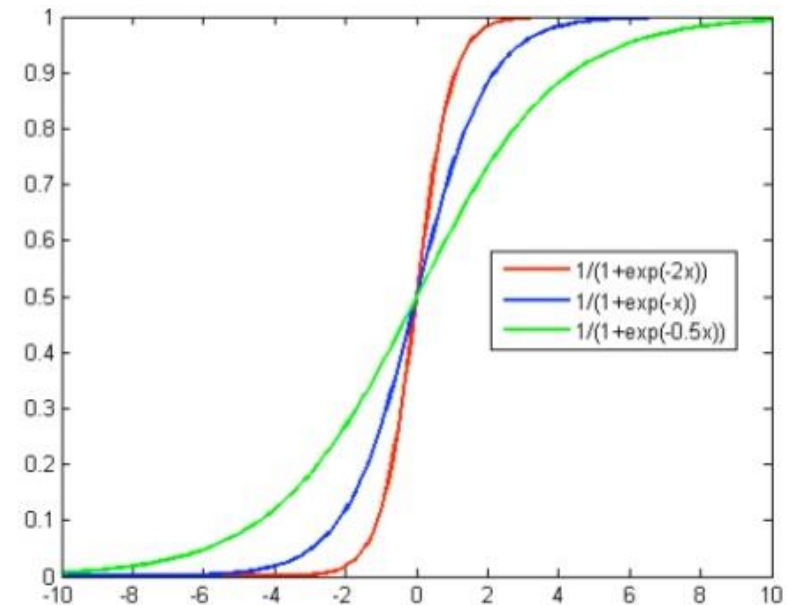
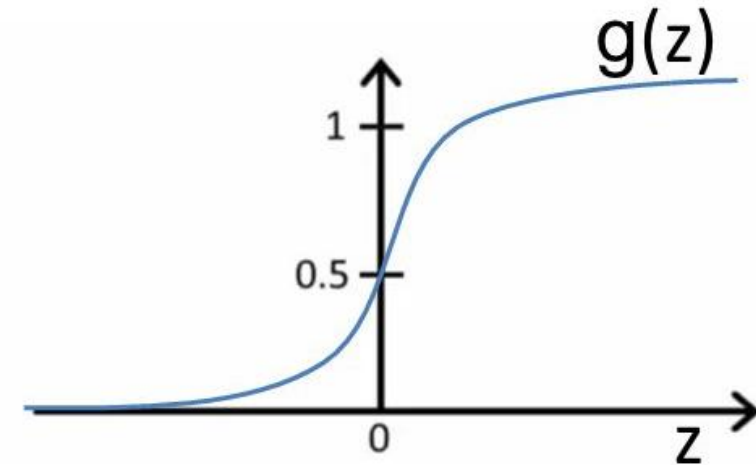
$$g(z) = \frac{1}{1 + e^{-z}}$$

- Define  $h_{\theta}(x) = g(\theta^T x)$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

- Predict parameters  $\theta$  to optimize  $h_{\theta}(x)$

But what does  $h_{\theta}(x)$  mean now??



# Logistic Regression - Meaning

- Interpretation of Hypothesis function

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

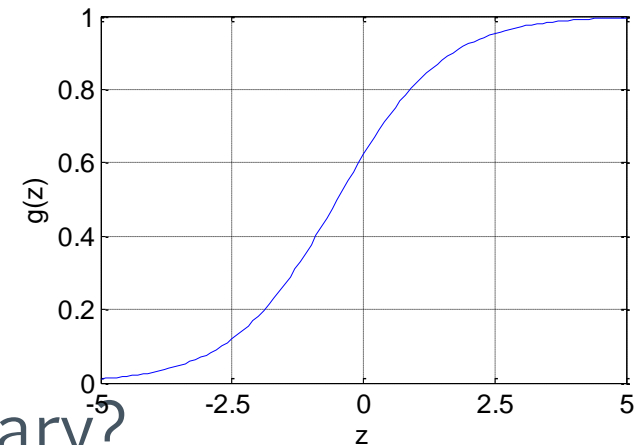
**$h_{\theta}(x)$**  = estimated probability that  $y = 1$  on input  $x$

Example:

$$h_{\theta}(x) = p(y = 1|x; \theta) = 0.4$$

Means that the probability that  $y=1$ , given  $x$ , parameterized by  $\theta$  is equal to 0.6

- Accordingly,  $p(y = 0|x; \theta) = 0.6$
- So the algorithm should predict  
 “ $y=1$ ” if  $h_{\theta}(x) \geq 0.6$  and “ $y=0$ ” if  $h_{\theta}(x) \leq 0.6$
- Which means also that  
 “ $y=1$ ” if  $z \geq 0$  and “ $y=0$ ” if  $z \leq 0$

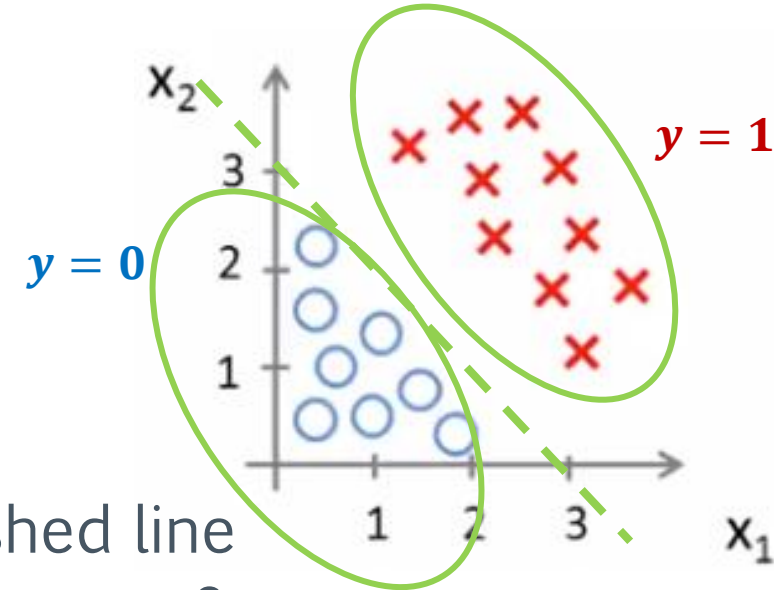


But how can we define this decision boundary?

# Logistic Regression-Decision Boundary

Consider the training set in the graph:

- For all points marked **x** the output should be  $y = 1$
- For all points marked **o** the output should be  $y = 0$



- **Decision Boundary** is shown as the dashed line
- Line equation:  $x_1 + x_2 = 3 \rightarrow -3 + x_1 + x_2 = 0$

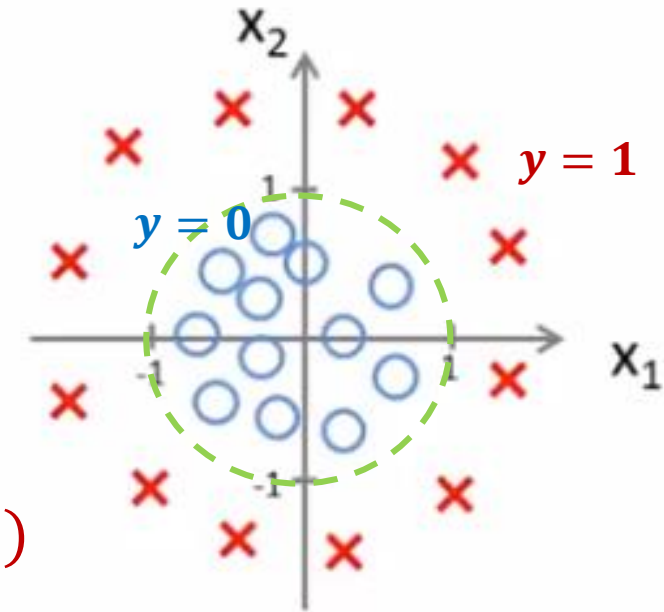
Thus

- Predict " $y = 1$ " if  $\rightarrow -3 + x_1 + x_2 \geq 0 : x_1 + x_2 \geq 3$
- Predict " $y = 0$ " if  $\rightarrow -3 + x_1 + x_2 < 0 : x_1 + x_2 < 3$

So  $\mathbf{h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2) = g(-3 + x_1 + x_2)}$

# Logistic Regression-Decision Boundary

- **Decision Boundaries** need not to be always linear, but could be any function that describes any shape that fits our data
- Consider the training set in the graph:



- We can predict hypothesis function as

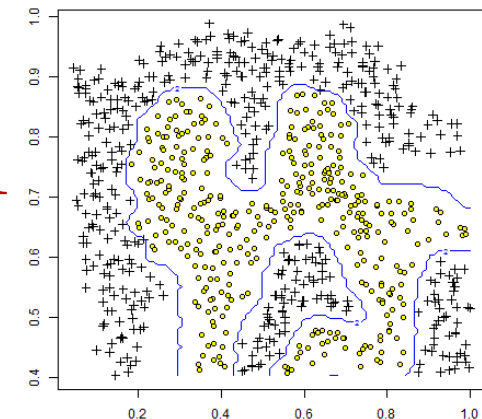
$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

- Using the training data we fit parameters  $\theta$

$$\theta^T = [-1 \quad 0 \quad 0 \quad 1 \quad 1]$$

- Predict " $y = 1$ " if  $\rightarrow -1 + x_1^2 + x_2^2 \geq 0 : x_1 + x_2 \geq 1$

- More complex hypotheses functions can be predicted to present decision boundaries for training data





# Logistic Regression – Cost Function

## ➤ Cost function in linear regression

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$y^i$  is  
continuous

Square  
the  
error

## ➤ Cost function for logistic regression

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

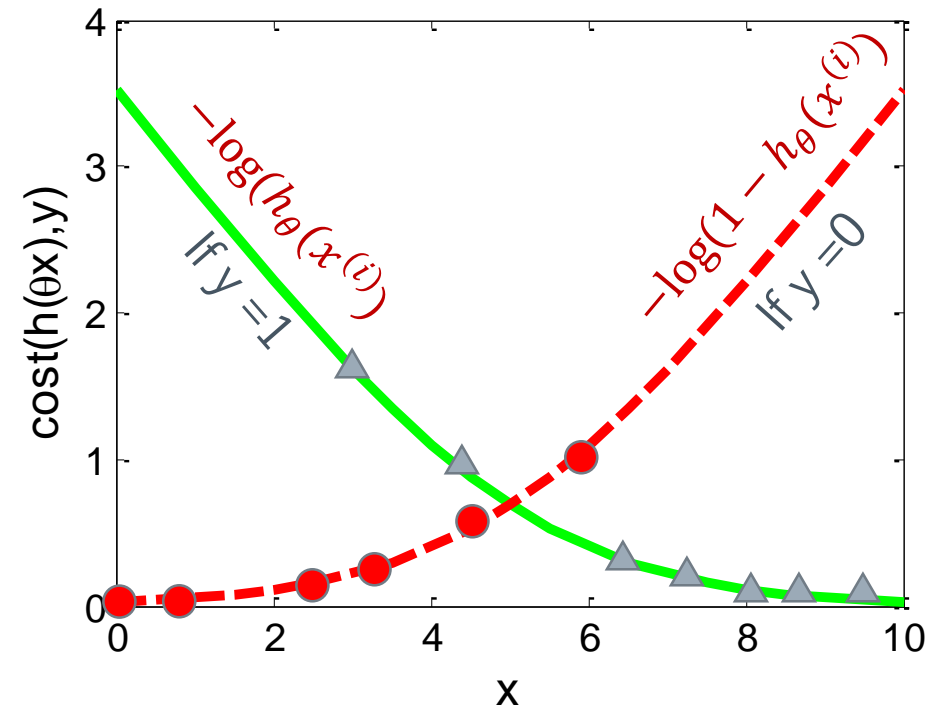
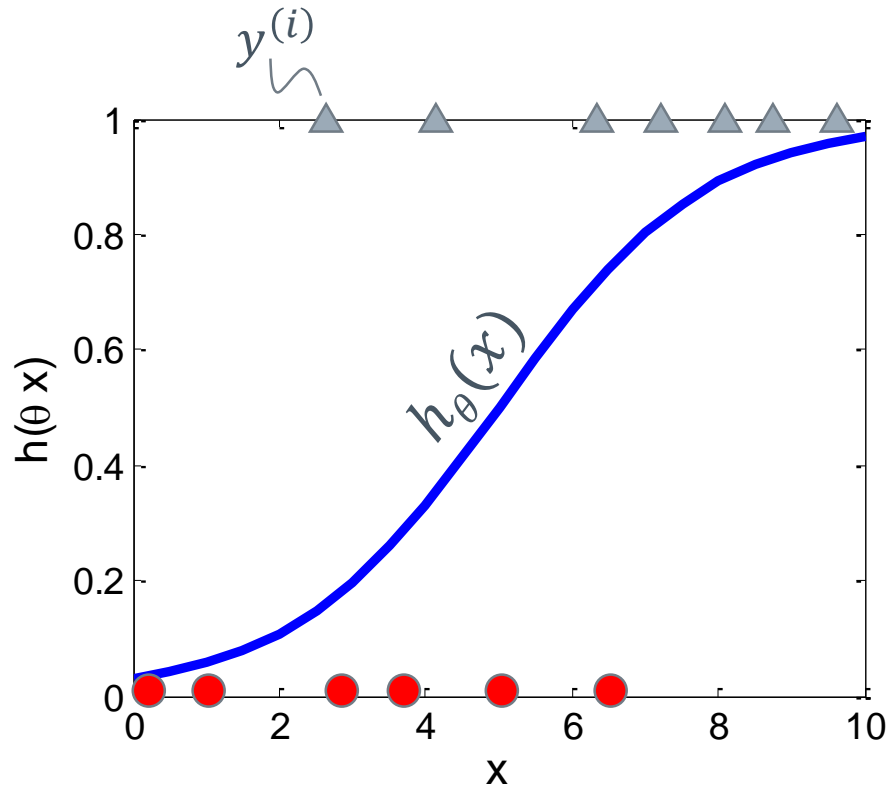
$y = 0$

$y = 1$

$$\text{Cost}(h_{\theta}(x^{(i)}), 1) = -\log(h_{\theta}(x^{(i)}))$$

$$\text{Cost}(h_{\theta}(x^{(i)}), 1) = -\log(1 - h_{\theta}(x^{(i)}))$$

# Logistic Regression Cost Function (Example)



# Logistic Regression Combined Cost Function

$$\text{Cost}(h_{\theta}(x^{(i)}), 1) = -\log(h_{\theta}(x^{(i)}))$$

$$\text{Cost}(h_{\theta}(x^{(i)}), 0) = -\log(1 - h_{\theta}(x^{(i)}))$$

$$\text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) = -y^{(i)}\log(h_{\theta}(x^{(i)})) - (1 - y^{(i)})\log(1 - h_{\theta}(x^{(i)}))$$

Will only have value  
when  $y^{(i)} = 1$



Will only have value  
when  $y^{(i)} = 0$

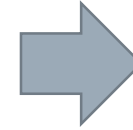
Average cost  $J(\theta)$  for a given Hypothesis

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m -y^{(i)}\log(h_{\theta}(x^{(i)})) - (1 - y^{(i)})\log(1 - h_{\theta}(x^{(i)}))$$

# Logistic Regression – Gradient Calculation

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m -y^{(i)} \log(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)}))$$

$$J(\theta) = \sum \frac{1}{m} \sum_{i=1}^m -y^{(i)} \log(a) - (1 - y^{(i)}) \log(1 - a)$$



$$a = h_{\theta}(x^{(i)}) = \frac{1}{1 + e^{-z}}$$

where  $z = \theta^T x$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{\partial J(\theta)}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial \theta_j}$$

# Logistic Regression – Gradient Calculation

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{\partial J(\theta)}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial \theta_j}$$

$$\begin{aligned} \frac{\partial J(\theta)}{\partial a} &= \frac{1}{m} \sum_{i=1}^m -y^{(i)} \frac{1}{a} + (1 - y^{(i)}) \frac{1}{1 - a} \\ &= \frac{1}{m} \sum_{i=1}^m \frac{-y^{(i)}(1 - a) + a(1 - y^{(i)})}{a(1 - a)} \\ &= \frac{1}{m} \sum_{i=1}^m \frac{a - y^{(i)}}{a(1 - a)} \end{aligned}$$

$$\begin{aligned} \frac{\partial a}{\partial z} &= \frac{-e^{-z}}{(1 + e^{-z})^2} \\ &= \frac{1}{1 + e^{-z}} \left( 1 - \frac{1}{1 + e^{-z}} \right) \\ &= a(1 - a) \end{aligned}$$

$$\frac{\partial z}{\partial \theta_j} = x_j^{(i)}$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m \frac{a - y^{(i)}}{a(1 - a)} a(1 - a) x_j^{(i)}$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$


# Logistic Regression – Gradient Descent

- › Similar to linear regression, gradient descent calculates  $\theta$  using the following iteration

$$\text{Repeat } \{ \theta_j = \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta) \text{ for } j=0,1,\dots,n \}$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m x_j^{(i)} [h_{\theta}(x^{(i)}) - y^{(i)}]$$

$$\text{Repeat } \{ \theta_j = \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m x_j^{(i)} [h_{\theta}(x^{(i)}) - y^{(i)}] \text{ for } j = 0,1, \dots n \}$$



$$\theta = \theta - \frac{\alpha}{m} X^T (H_{\theta}(X) - Y)$$

# Logistic Regression Gradient Descent - Example

$$X = \begin{bmatrix} 1 & 2 & 4 \\ 1 & 1 & 7 \\ 1 & 5 & 10 \\ 1 & 6 & 2 \end{bmatrix} \quad Y = \begin{bmatrix} 1 \\ 1 \\ 0 \\ 0 \end{bmatrix} \quad \theta = \begin{bmatrix} -4 \\ -1 \\ 1 \end{bmatrix} \quad \alpha = 0.3$$

$$H_{\theta}(X) = \frac{1}{1 + \exp(-X\theta)}$$

$$= \frac{1}{1 + \exp\left(-\begin{bmatrix} 1 & 2 & 4 \\ 1 & 1 & 7 \\ 1 & 5 & 10 \\ 1 & 6 & 2 \end{bmatrix} \begin{bmatrix} -4 \\ -1 \\ 1 \end{bmatrix}\right)}$$

$$H_{\theta}(X) = \begin{bmatrix} 0.1192 \\ 0.8808 \\ 0.7311 \\ 0.0003 \end{bmatrix}$$

$$\theta_{new} = \begin{bmatrix} -4 \\ -1 \\ 1 \end{bmatrix} - \frac{0.3}{4} \begin{bmatrix} 1 & 2 & 4 \\ 1 & 1 & 7 \\ 1 & 5 & 10 \\ 1 & 6 & 2 \end{bmatrix}^T \left( \begin{bmatrix} 0.1192 \\ 0.8808 \\ 0.7311 \\ 0.0003 \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \\ 0 \\ 0 \end{bmatrix} \right)$$

$$\theta_{new} = \begin{bmatrix} -3.9799 \\ -1.1332 \\ 0.7785 \end{bmatrix}$$

$$\theta = \theta - \frac{\alpha}{m} X^T (H_{\theta}(X) - Y)$$

# Machine Learning Fundamentals – DTSC102

## Lecture 6 ML Diagnostics

Course Instructor: Dr.-Ing. Maggie Mashaly  
maggie.ezzat@guc.edu.eg  
C3.220



# Contents

- Machine Learning Diagnostics
- Evaluating your Hypothesis
- Model Selection
- Diagnosing Bias vs. Variance

# Machine Learning Diagnostic

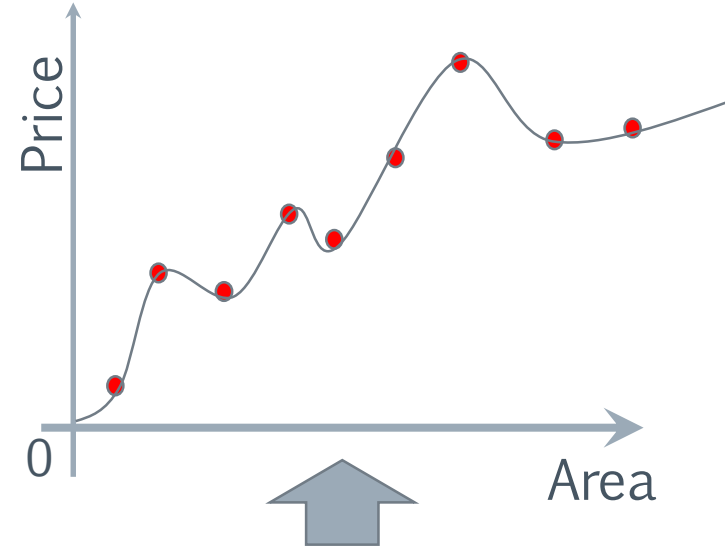
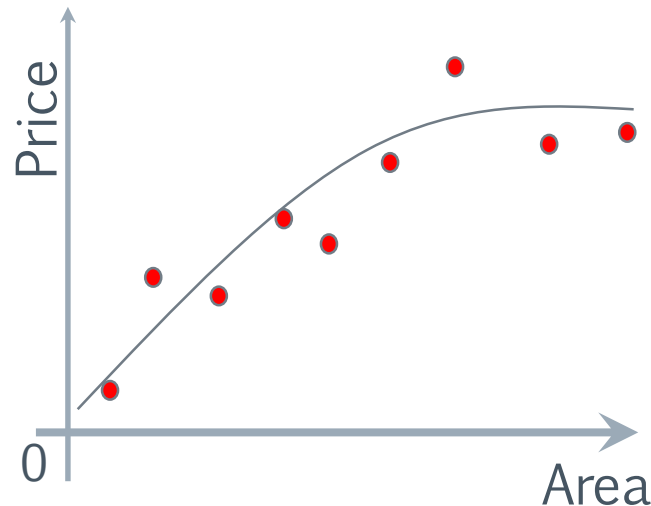
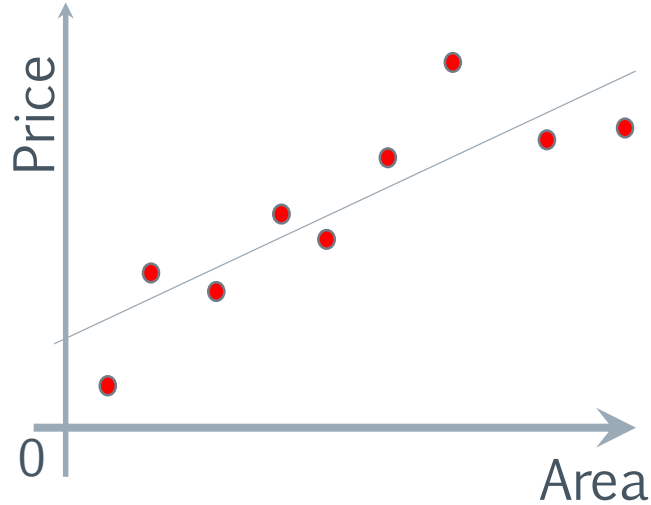
- Suppose you have implemented linear regression to predict housing prices.

$$J(\theta_0, \theta_1, \dots, \theta_m) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- However, when you test your hypothesis on a new set of houses, you find that it makes unacceptably large errors in its predictions. What should you try next?
  - Get more training examples
  - Try adding polynomial features  $(x_1^2, x_2^2, x_1 x_2)$
  - Try smaller sets of features
  - Try a different hypothesis
  - Try getting additional features

# Choosing your Hypothesis

Which one is the best hypothesis ??



This one will present the lowest cost function (lowest error)

But this may not present a general learning

It represents an **Over Fitting** of the data

# Over Fitting & Under Fitting

## Over Fitting

- › Excellent fitting of the training
- › Poor ability to predict unlearned data
- › Called **high variance**
- › Caused by
  - Very complex Hypothesis
  - Too many features

## Under Fitting

- › Poor Fitting of the learning data
- › Called **high bias**
- › Caused by:
  - Simple hypothesis
  - Too few features
  - Small number of training data

## Solution

Model Selection

Regularization

# Model Selection

Consider your given dataset:

- Split the data into a **Training Set** (70%) and a **Test Set** (30%)

For **Linear** Regression:

- Learn parameter  $\theta$  from training data to minimize training error  $J(\theta)$
- Compute test set error:

$$J_{test}(\theta) = \frac{1}{2m_{test}} \sum_{i=1}^{m_{test}} \left( h_{\theta} \left( x_{test}^{(i)} \right) - y_{test}^{(i)} \right)^2$$

	Size	Price
Training Set Size= $m$	2104	400
	1600	330
	2400	369
	1416	232
	3000	540
	1985	300
Test Set Size= $m_{test}$	1534	315
	1427	199
	1380	212
	1494	243

# Model Selection

For **Logistic** Regression:

➤ Learn parameter  $\theta$  from training data to minimize training error  $J(\theta)$

➤ Compute test set error:

$$J_{test}(\theta) = -\frac{1}{m_{test}} \sum_{i=1}^{m_{test}} y_{test}^{(i)} \log h_{\theta}(x_{test}^{(i)}) + (1 - y_{test}^{(i)}) \log h_{\theta}(x_{test}^{(i)})$$

➤ Alternatively, we can calculate **Misclassification Error (0/1 Error)**

$$err(h_{\theta}(x), y) = \begin{cases} 1 & \text{if } h_{\theta}(x) \geq 0.5 \text{ and } y = 0 \text{ or } h_{\theta}(x) < 0.5 \text{ and } y = 1 \\ 0 & \text{Otherwise} \end{cases}$$

which gives us a binary 0/1 error based on a misclassification, then calculate **Average Test error**

$$Test\ Error = \frac{1}{m_{test}} \sum_{i=1}^{m_{test}} err(h_{\theta}(x_{test}^{(i)}) - y_{test}^{(i)})$$

# Model Selection

Which polynomial to choose for Hypothesis Function??

$$1. h_{\theta}(x) = \theta_0 + \theta_1 x \quad \text{for } d=1 : \theta^{(1)} \rightarrow J_{test}(\theta^{(1)})$$

$$2. h_{\theta}(x) = \theta_0 + \theta_1 x + \theta_2 x^2 \quad \text{for } d=2 : \theta^{(2)} \rightarrow J_{test}(\theta^{(2)})$$

$$3. h_{\theta}(x) = \theta_0 + \theta_1 x + \dots + \theta_3 x^3 \quad \text{for } d=3 : \theta^{(3)} \rightarrow J_{test}(\theta^{(3)})$$

⋮

$$10. h_{\theta}(x) = \theta_0 + \theta_1 x + \dots + \theta_{10} x^{10} \quad \text{for } d=10 : \theta^{(10)} \rightarrow J_{test}(\theta^{(10)})$$

➤ Choose degree with the lowest test-set error (ex:  $h_{\theta}(x) = \theta_0 + \dots + \theta_5 x^5$ )

- But does this model generalize well? Not really...

- Why? Because  $J_{test}(\theta^{(5)})$  is likely to be an optimistic estimate of generalization error, as  $d$  was chosen to fit the test set

# Model Selection

Back to our dataset:

- Split the data into a **Training Set** (60%), a **Cross Validation (CV) Set** (20%) and a **Test Set** (20%)
- Optimize the parameters in  $\theta$  using the training set for each polynomial degree  $J_{train}(\theta)$
- Find the polynomial degree  $d$  with the least error using the cross validation set.

$$J_{cv}(\theta) = \frac{1}{2m_{cv}} \sum_{i=1}^{m_{cv}} \left( h_{\theta} \left( x_{cv}^{(i)} \right) - y_{cv}^{(i)} \right)^2$$

- Estimate the generalization error using the test

$$\text{set } J_{test}(\theta) = \frac{1}{2m_{test}} \sum_{i=1}^{m_{test}} \left( h_{\theta} \left( x_{test}^{(i)} \right) - y_{test}^{(i)} \right)^2$$

	Size	Price
Training Set Size= $m$	2104	400
	1600	330
	2400	369
	1416	232
	3000	540
Cross Validation Set Size= $m_{cv}$	1985	300
	1534	315
	1427	199
Test Set Size= $m_{test}$	1380	212
	1494	243



# Model Selection

Which polynomial to choose for Hypothesis Function??

- |  |  |
|--|--|
| 1. $h_{\theta}(x) = \theta_0 + \theta_1 x$                               | for $d=1 : \theta^{(1)} \rightarrow J_{cv}(\theta^{(1)})$    |
| 2. $h_{\theta}(x) = \theta_0 + \theta_1 x + \theta_2 x^2$                | for $d=2 : \theta^{(2)} \rightarrow J_{cv}(\theta^{(2)})$    |
| 3. $h_{\theta}(x) = \theta_0 + \theta_1 x + \dots + \theta_3 x^3$        | for $d=3 : \theta^{(3)} \rightarrow J_{cv}(\theta^{(3)})$    |
| ⋮  |  |
| 10. $h_{\theta}(x) = \theta_0 + \theta_1 x + \dots + \theta_{10} x^{10}$ | for $d=10 : \theta^{(10)} \rightarrow J_{cv}(\theta^{(10)})$ |

- Choose degree with the lowest cross- validation set error  
 (ex:  $h_{\theta}(x) = \theta_0 + \dots + \theta_5 x^4$ )
- Estimate generalization error for test set  $J_{test}(\theta^{(4)})$

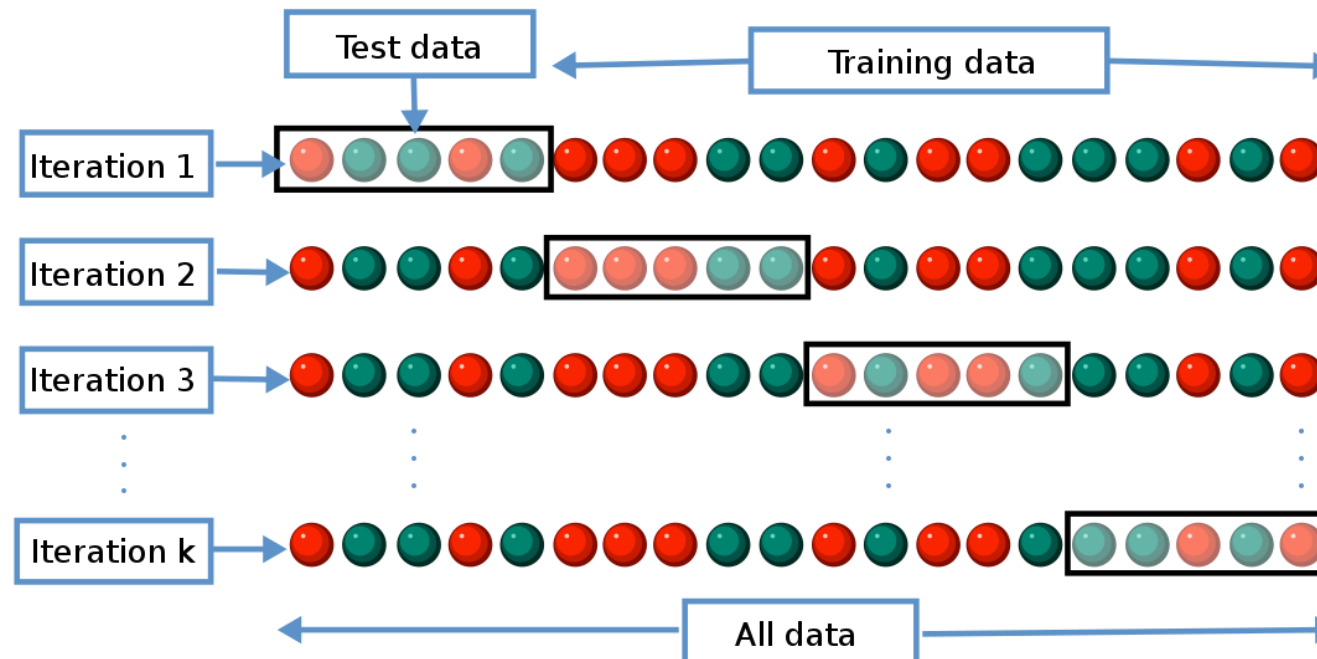
(Thus the degree of the polynomial  $d$  has not been trained using the test set)

# Model Selection

However there are still limitations for using a single training/CV/test set:

- Original data might not be enough to make a sufficiently large training/CV/test sets
- A single training set doesn't tell us how sensitive accuracy is to a particular training sample

Possible solution: **Random Resampling**

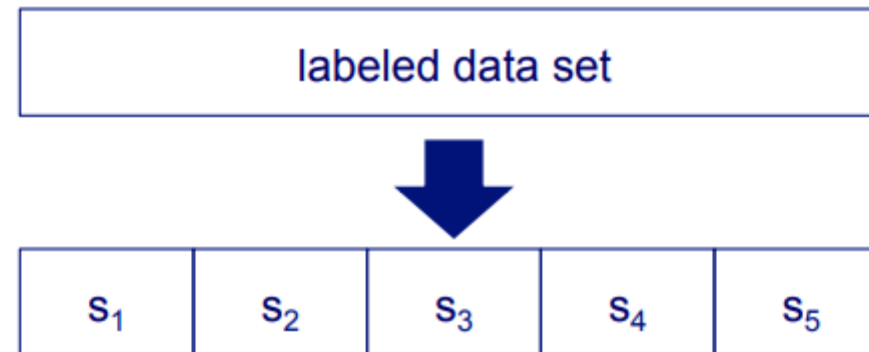


# Model Selection

Done more accurately it is called: **K-fold Sampling**

- Partition data into n subsamples
- Iteratively leave one subsample out for the test set, train on the rest

- Suppose we have 100 instances:  
 $\text{Accuracy} = 73/100 = 73\%$
- Common value for K is 10, smaller numbers are also used when learning is time consuming



iteration	train on	test on	correct
1	S <sub>2</sub> S <sub>3</sub> S <sub>4</sub> S <sub>5</sub>	S <sub>1</sub>	11 / 20
2	S <sub>1</sub> S <sub>3</sub> S <sub>4</sub> S <sub>5</sub>	S <sub>2</sub>	17 / 20
3	S <sub>1</sub> S <sub>2</sub> S <sub>4</sub> S <sub>5</sub>	S <sub>3</sub>	16 / 20
4	S <sub>1</sub> S <sub>2</sub> S <sub>3</sub> S <sub>5</sub>	S <sub>4</sub>	13 / 20
5	S <sub>1</sub> S <sub>2</sub> S <sub>3</sub> S <sub>4</sub>	S <sub>5</sub>	16 / 20

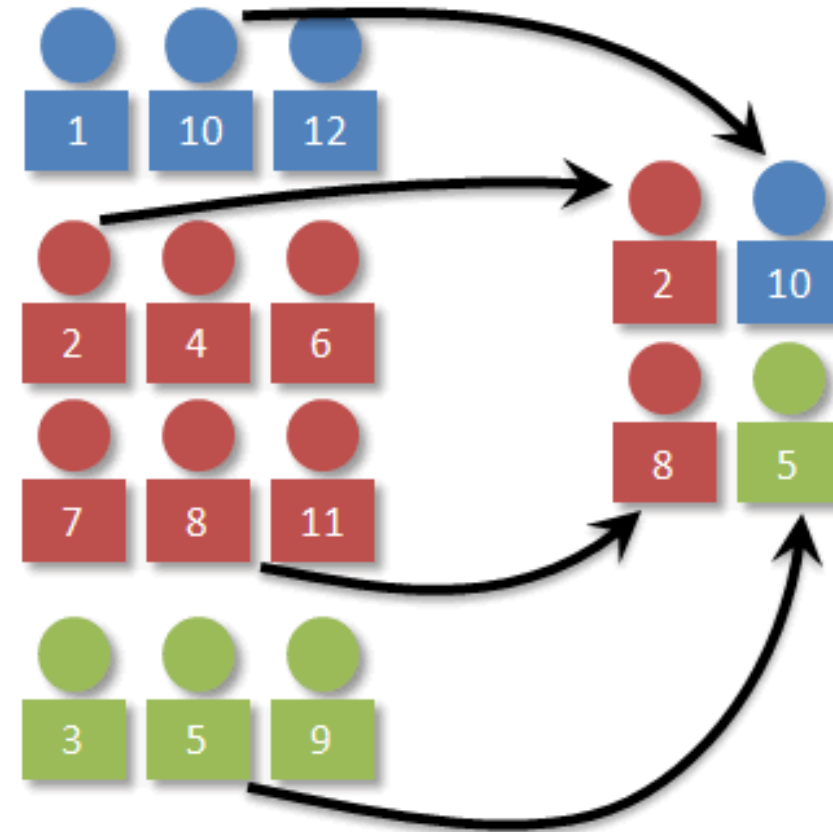
# Model Selection

Other possible solution: **Stratified Sampling**

➤ Ensures that class proportions are maintained in each selected set

➤ How it is done:

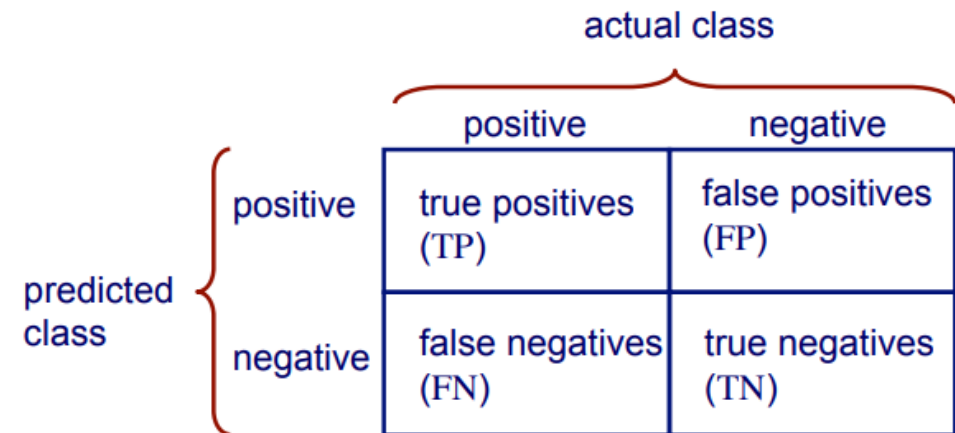
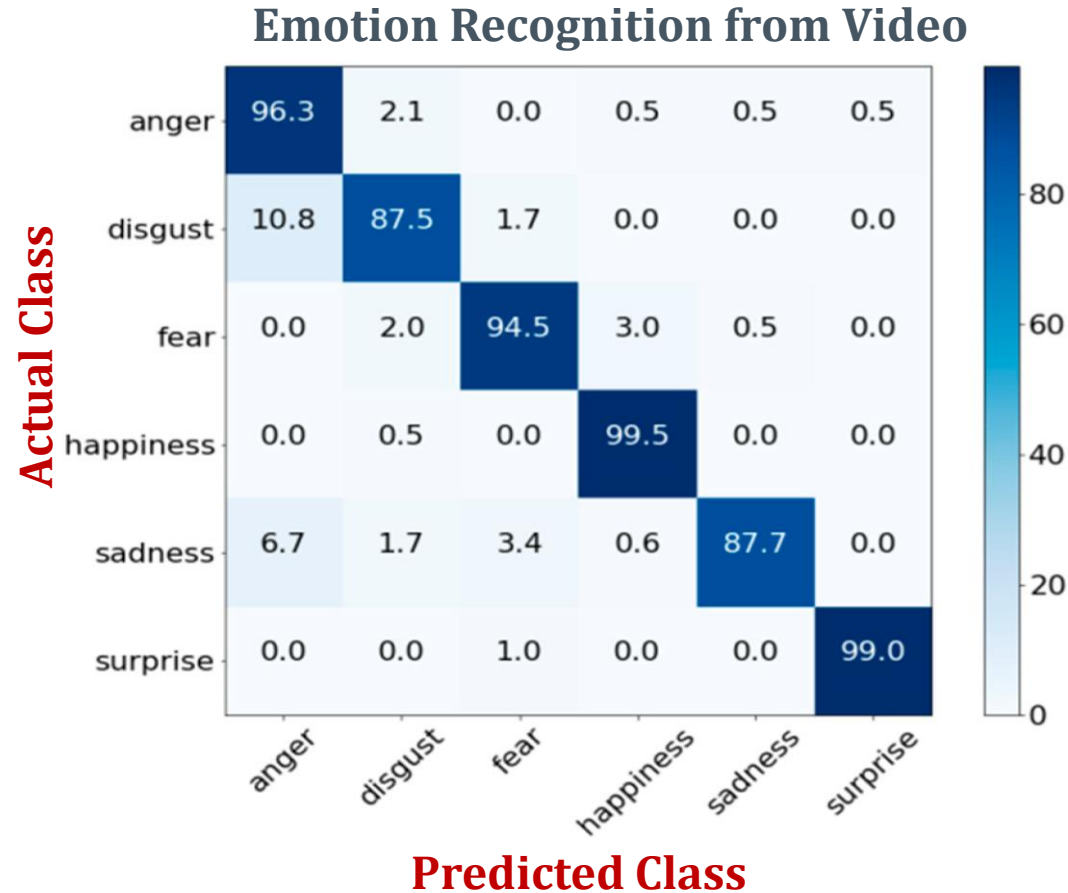
1. Stratify instances by class
2. Randomly select instances from each class proportionally



# Model Evaluation

How to understand types of mistakes your model is making?

➤ Construct a **Confusion Matrix**



$$\text{accuracy} = \frac{TP + TN}{TP + FP + FN + TN}$$

# Model Evaluation

## Confusion Matrix

Is the simple accuracy calculation  $\frac{TP+TN}{TP+FP+FN+TN}$  sufficient? **Not Always...**

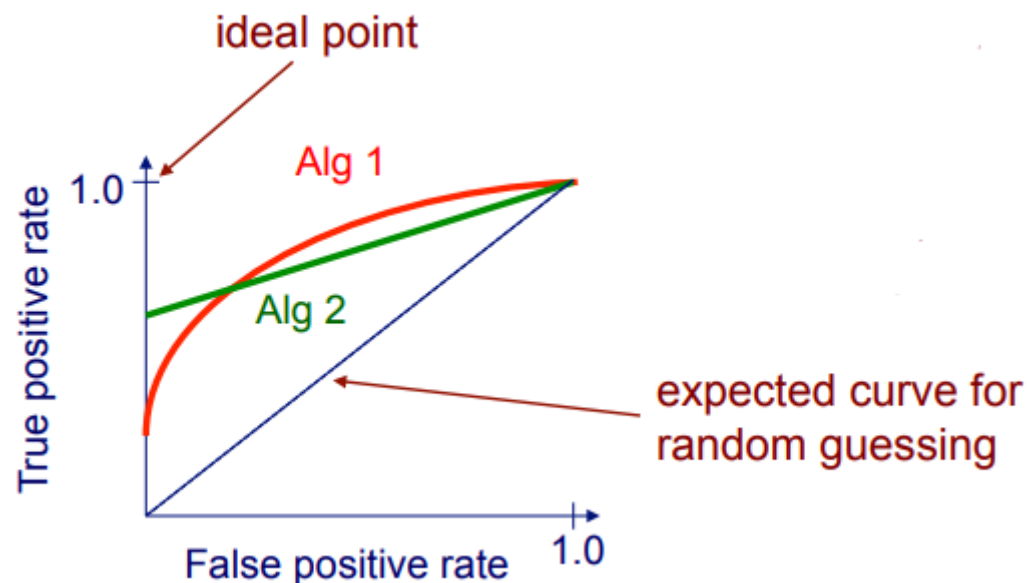
- Accuracy may not be useful measure in cases where there is a large class skew
  - Is 98% accuracy good if 97% of the instances are negative?
- There are differential misclassification costs; say, getting a positive wrong costs more than getting a negative wrong
  - Consider a medical domain in which a false positive results in an extraneous test but a false negative results in a failure to treat a disease
- We are most interested in a subset of high-confidence predictions

# Model Evaluation

## Confusion Matrix

➤ Other Accuracy metrics: TP-rate & FP-rate

➤ **ROC Curves: “Receiver Operating Characteristics”** curve plots the TP-rate (sensitivity/prob. Of detection) vs. the FP-rate (prob. of false alarm) at various threshold settings



		actual class	
		positive	negative
predicted class	positive	true positives (TP)	false positives (FP)
	negative	false negatives (FN)	true negatives (TN)

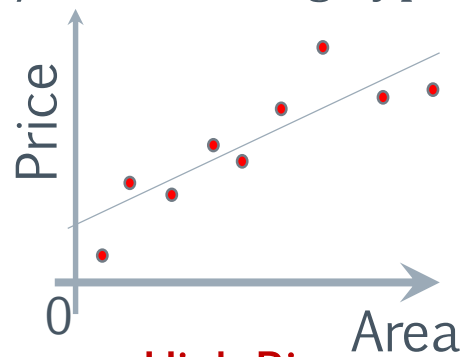
$$\text{true positive rate (recall)} = \frac{\text{TP}}{\text{actual pos}} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$
  

$$\text{false positive rate} = \frac{\text{FP}}{\text{actual neg}} = \frac{\text{FP}}{\text{TN} + \text{FP}}$$

# Model Evaluation

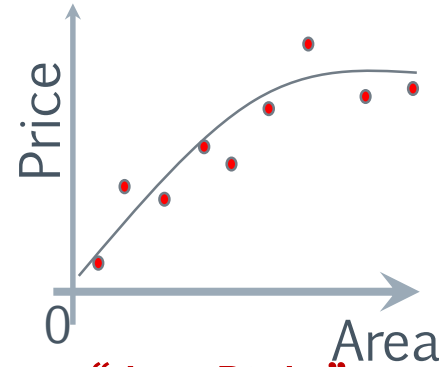
## Diagnosing Bias vs. Variance

Recall Over/Under fitting hypotheses

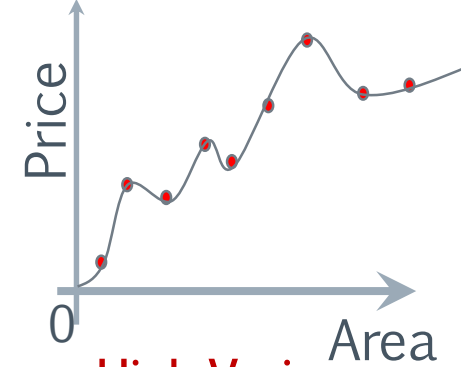


**High Bias  
(Underfit)**

**d=1**



**“Just Right”  
d=2**



**High Variance  
(Overfit)**

**d=4**

Training Error:

$$J_{Train}(\theta) = \frac{1}{2m_{Train}} \sum_{i=1}^{m_{Train}} \left( h_{\theta} \left( x_{Train}^{(i)} \right) - y_{Train}^{(i)} \right)^2$$

Cross Validation Error:

$$J_{cv}(\theta) = \frac{1}{2m_{cv}} \sum_{i=1}^{m_{cv}} \left( h_{\theta} \left( x_{cv}^{(i)} \right) - y_{cv}^{(i)} \right)^2$$



# Model Evaluation

## Diagnosing Bias vs. Variance

How to diagnose if your learning algorithm is suffering from a Bias/Variance Problem?

### Bias (Underfit)

- $J_{Train}(\theta)$  will be high
- $J_{CV}(\theta) \approx J_{Train}(\theta)$

### Variance (Overfit)

- $J_{Train}(\theta)$  will be low
- $J_{CV}(\theta) \gg J_{Train}(\theta)$

